

junia ISEN

Artificial Intelligence

Machine Learning (1)

Junia-ISEN / M1 / 2024-2025
Nacim Ihaddadene

Machine Learning VS classical programming

Let's consider the prototyping of the “Mean” function in Python:

```
def mean(x, y):  
    return (x + y) / 2  
  
mean_value = mean(3, 5)  
print(mean_value)
```

Let's do it now with ML !

Machine Learning VS classical programming

```
import numpy as np
from sklearn.linear_model import LinearRegression

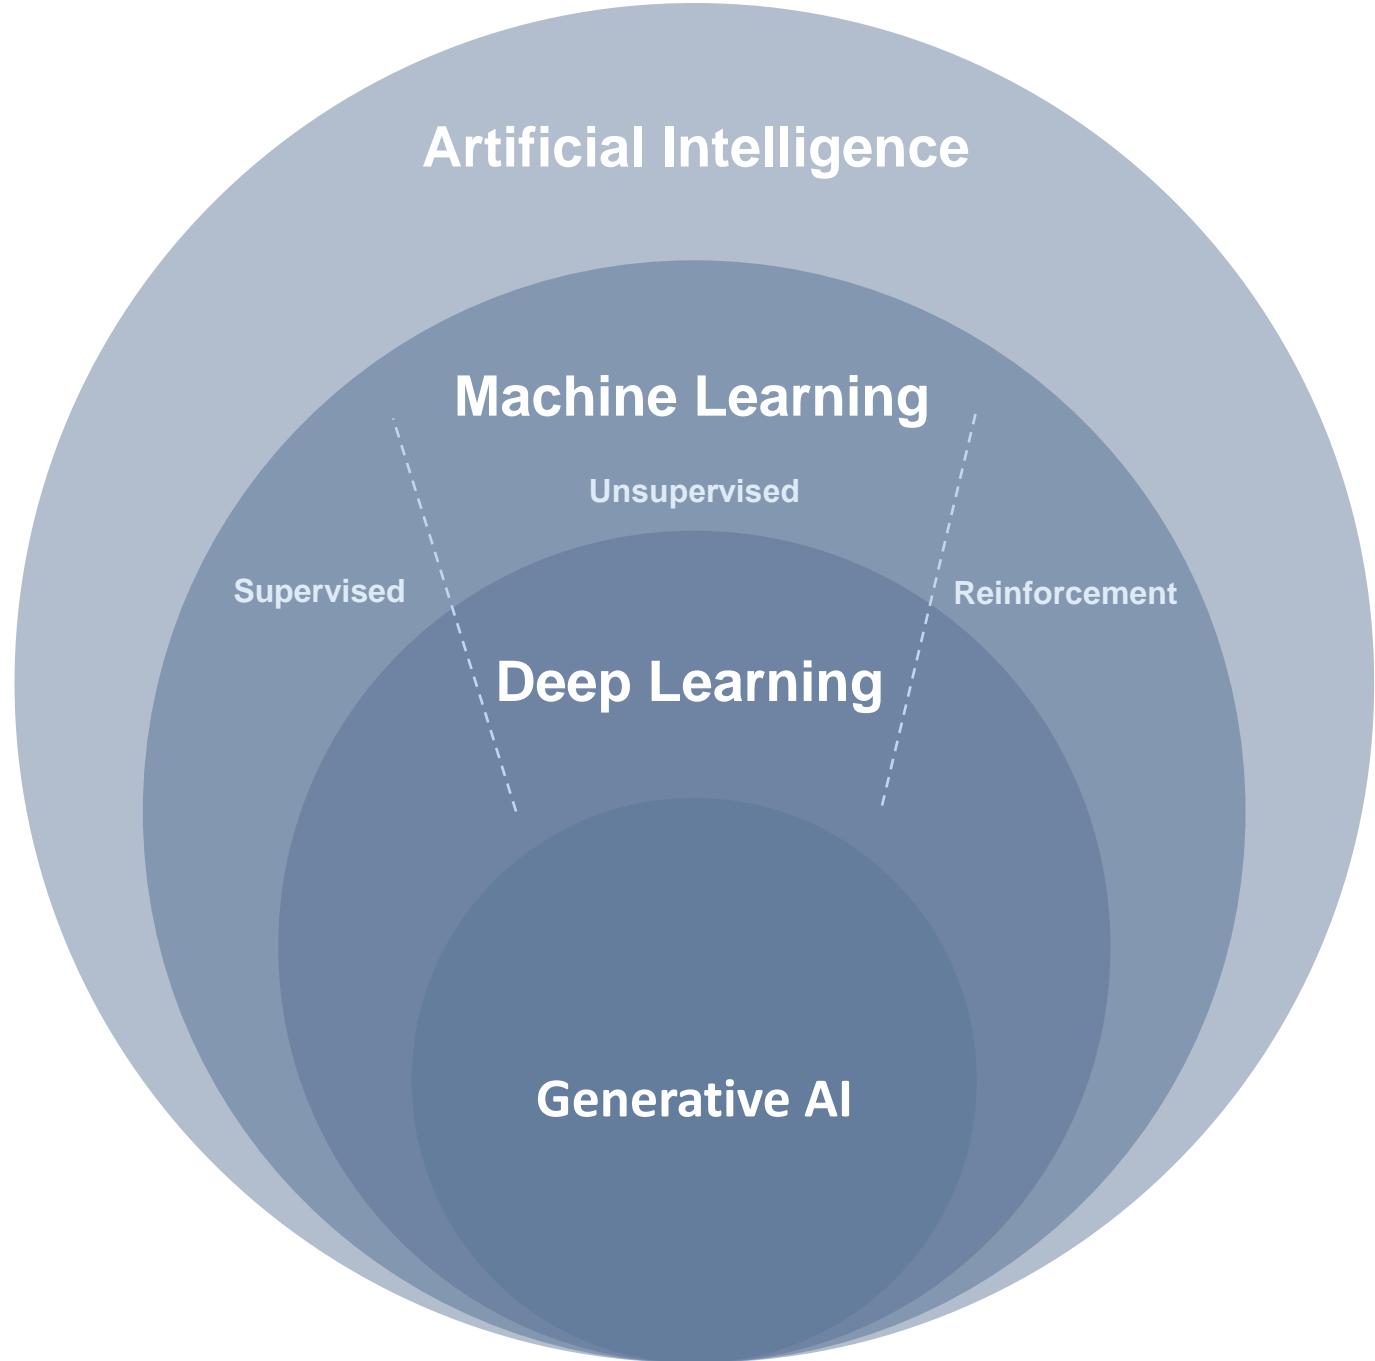
X = np.array([[2, 4], // array of input examples
              [5, 5],
              [7, 3],
              [2, 8],
              [7, 1]])
Y = np.array([3, 5, 5, 5, 4]) // array of corresponding outputs

model = LinearRegression()
model.fit(X, Y) // train the model
print(model.predict([[4, 6]])) // use it for a new prediction
```



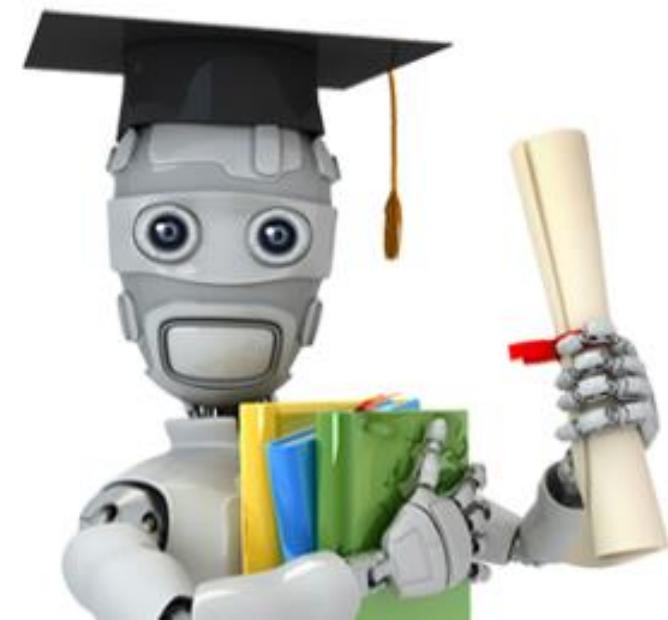
This example only shows the difference between the ways of doing things. Applying ML to calculate a mean value is not the best thing to do.

ML becomes interesting when you are overwhelmed by the number of parameters in the problem, and you have enough examples to model it.

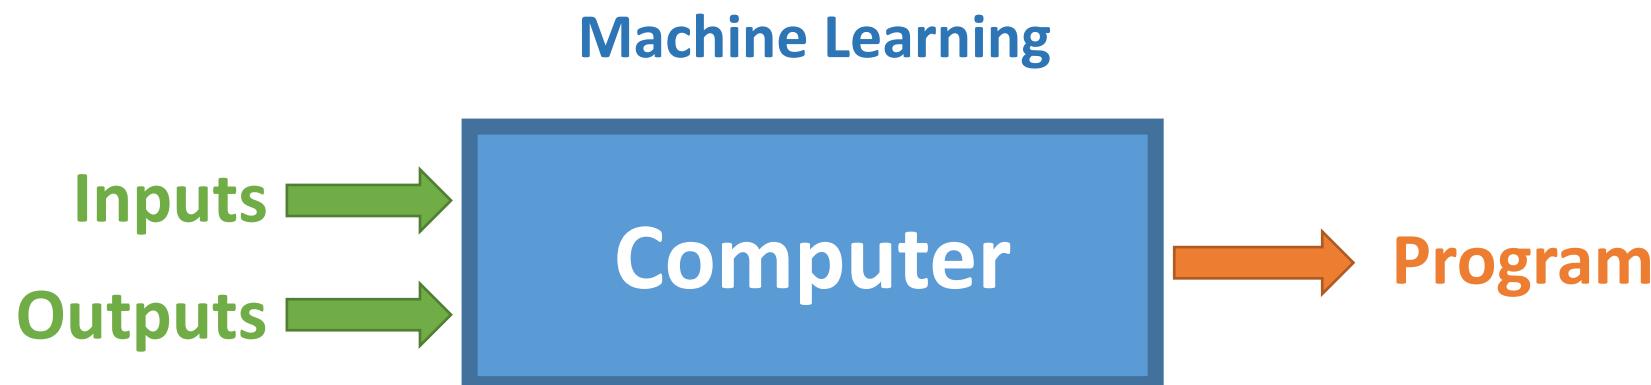


Machine Learning

- Giving computers the ability to learn **without being explicitly programmed.** Arthur Samuel
- Designing algorithms that can learn **patterns** from **data** (and exploit them)
- A subset of A.I. where the machine is trained to learn from it's **past experience**. The past experience is developed through the data collected.



ML vs Traditionnal Programming



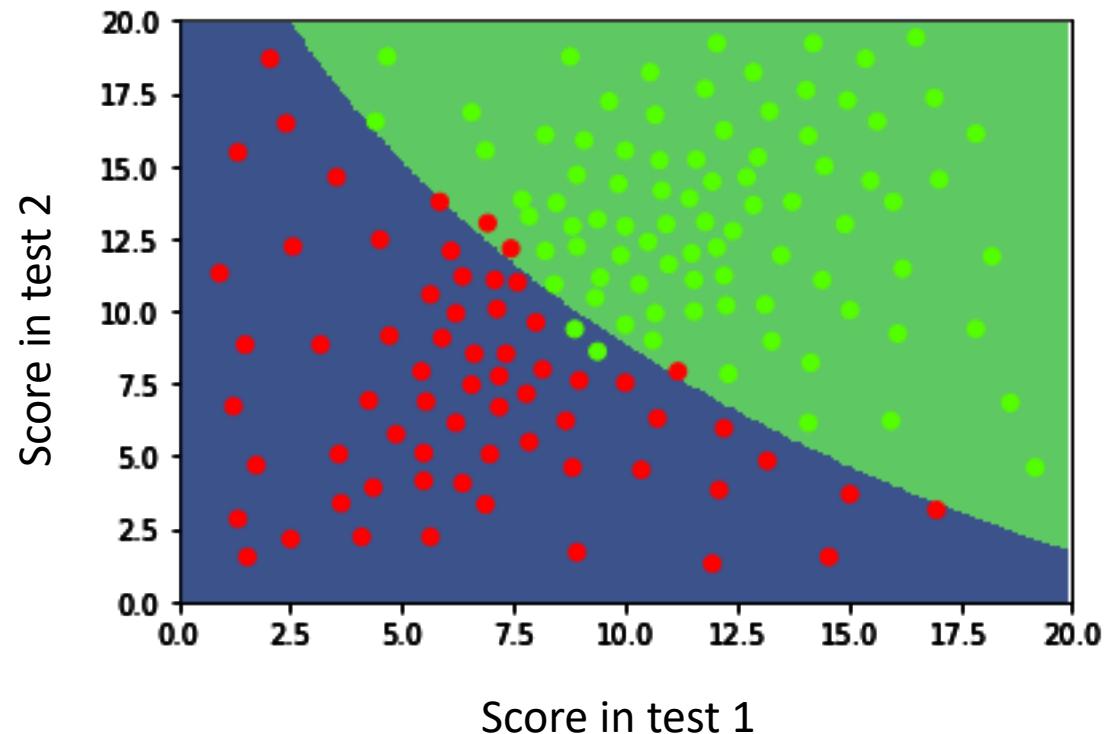
Example : College admissions

- Traditionnal programming :

```
If ((test1 + test2) / 2 >= 10)  
then : Accept  
Else : Reject
```

- Machine learning

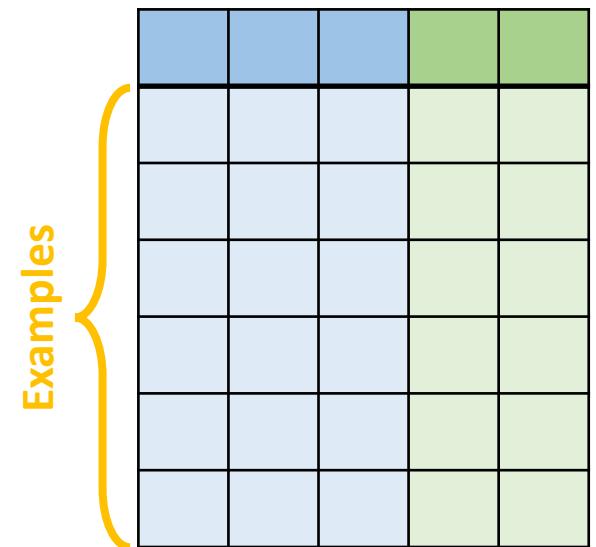
Given the history of acceptance or rejection of previous students, find a law that estimates the decision automatically for the future students



How to prepare **inputs** and **outputs**

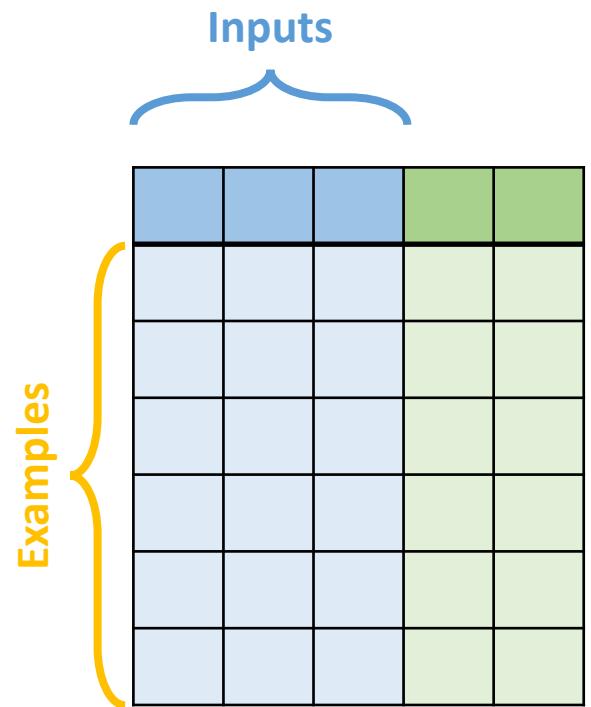
Machine Learning \leftrightarrow Learning by **examples**

- What's an **example**?
- We can talk about :
 - **Instance**, specific type of example
 - **Thing** or **element** to be classified, associated, or clustered
 - **Individual**, independent example of target concept
 - **Record**, characterized by a predetermined set of attributes
 - **Cases**,
- A **Dataset** is a big set of examples



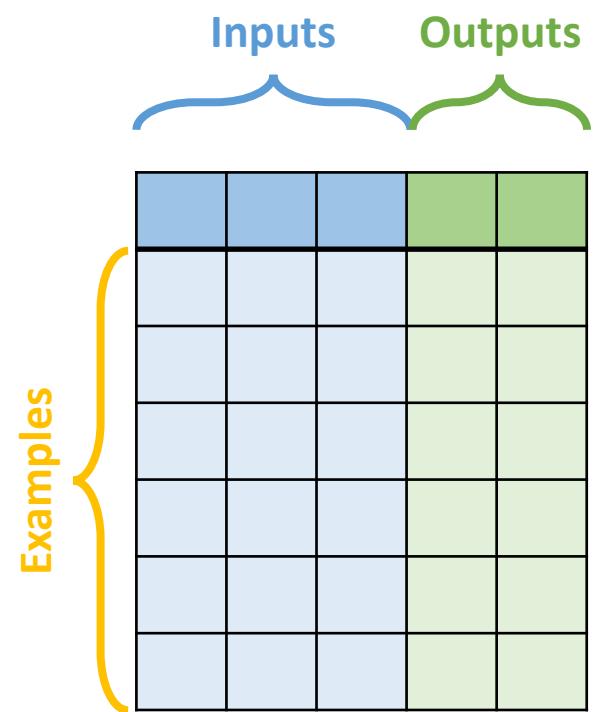
What's an **attribute**?

- Each Example or instance is described by a predefined set of **features** (also called **attributes**, or **fields**, or **characteristics**, ...)
- **Nominal values** : distinct and discrete
 - Name, Position, ...
 - Color : Blue, White, Red ...
- **Ordinal values**: could be ordered
 - Grade : A,B,C,D,E,...
 - Temperature : Cold, Mild, Hot
- **Continuous values**:
 - Height, Weight, Temperature in degrees, colors in RGB values



What's a label?

- A label is the desired **output** of the ML system
 - Possible synonyms : **category**, **answer**, **result**, ...
 - The labels corresponding to a dataset examples may be available or not, depending on the problem.

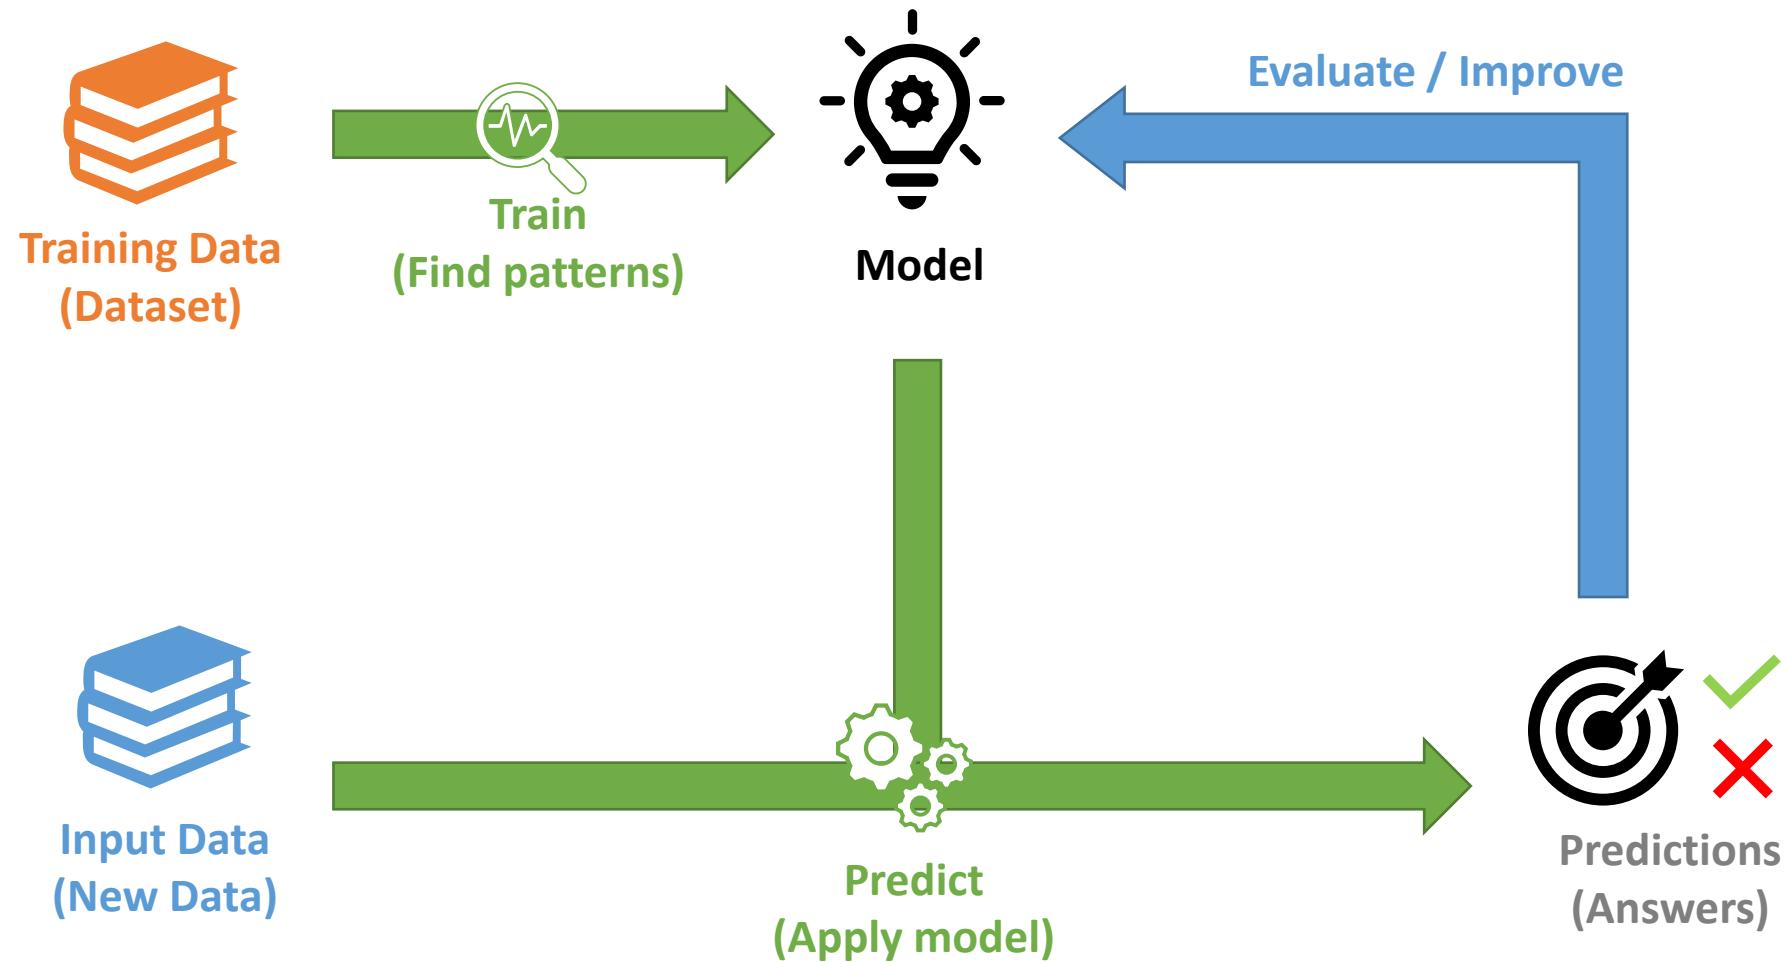


Back to college admission problem

- Each **example** is a candidate with his scores and decision
- The **attributes** are the scores in the two exams
- The **label** is the decision (accepted or rejected)
- The **dataset** is a long list of all previous candidate scores and decisions

	Score 1	Score 2	Decision
Student 1	17	14	Accepted
Student 2	7	9	Rejected
Student 3	10	13	Accepted
Student 4	9	15	Accepted
Student 5	11	2	Rejected
Student	19	1,5	Rejected
...			

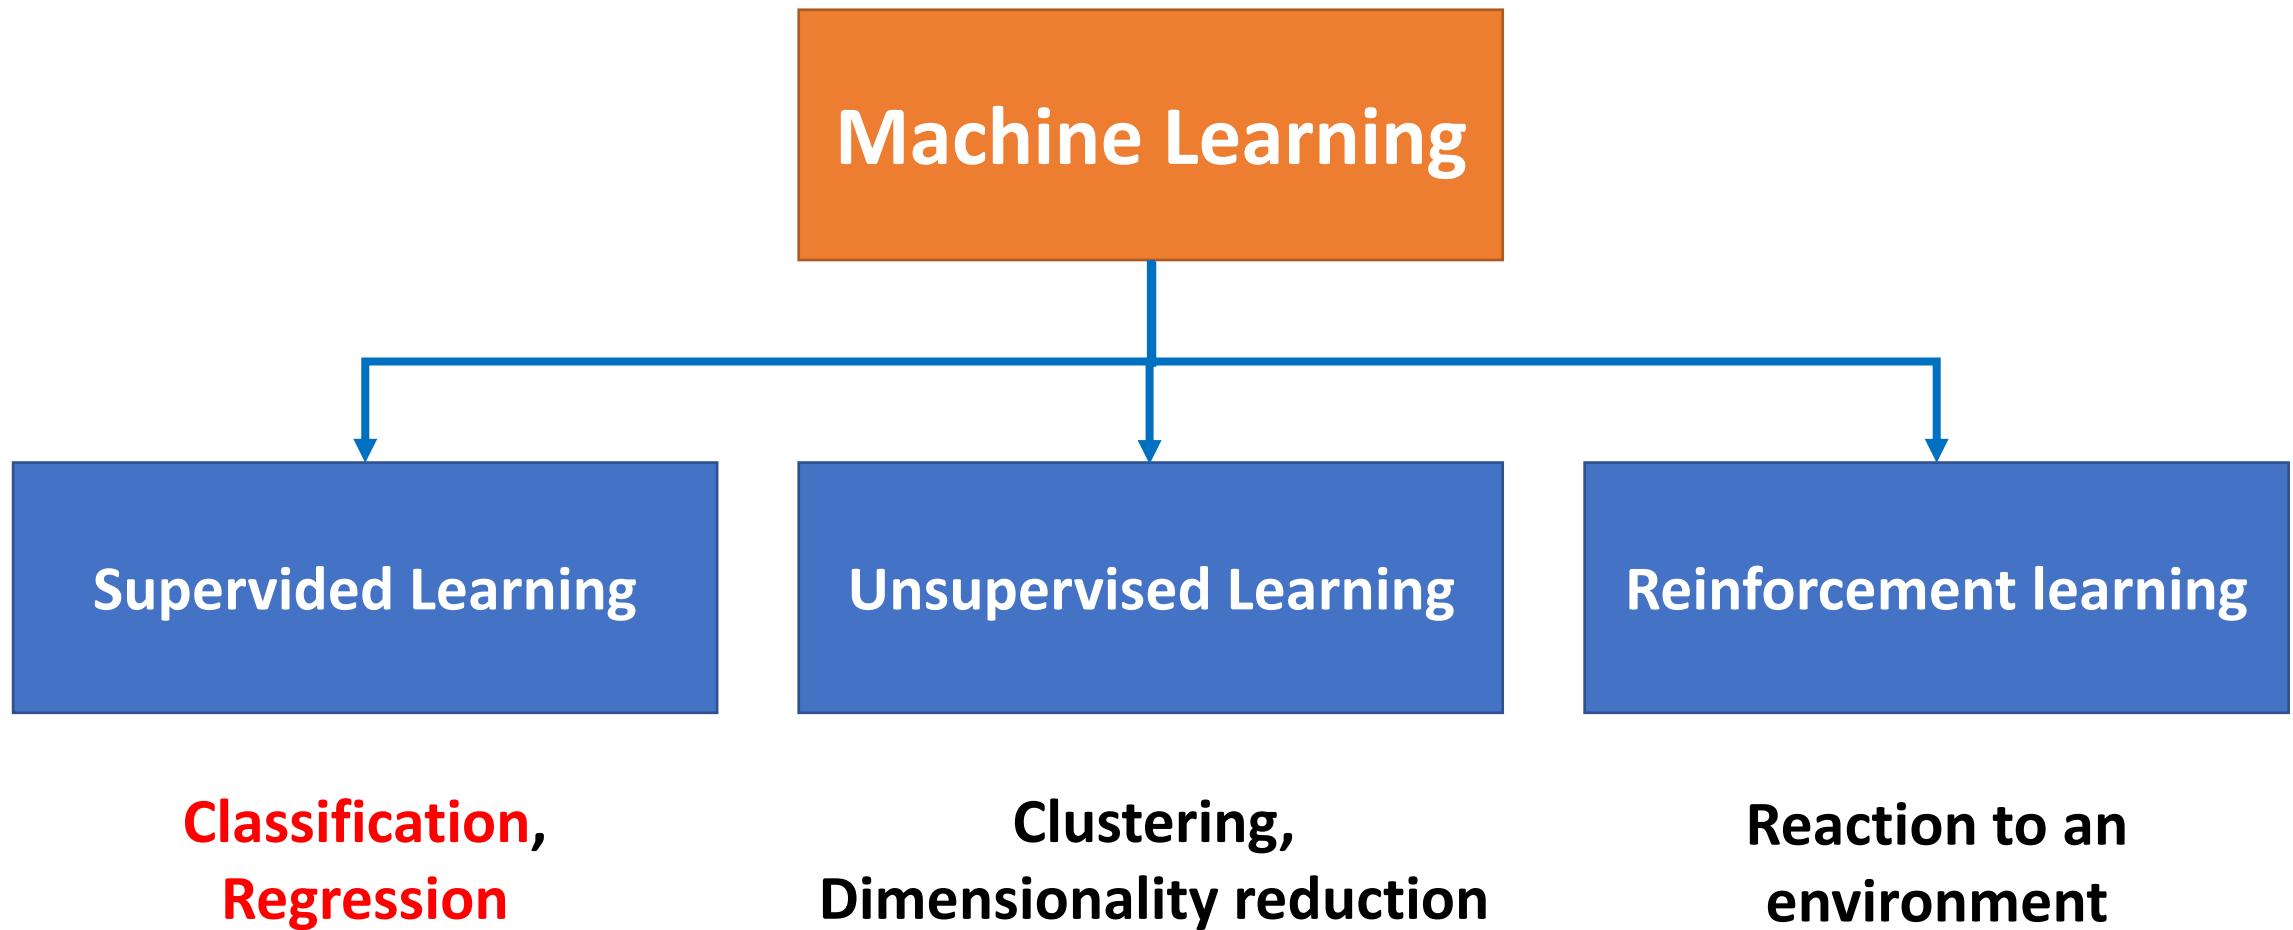
Machine Learning Process



Machine Learning Process

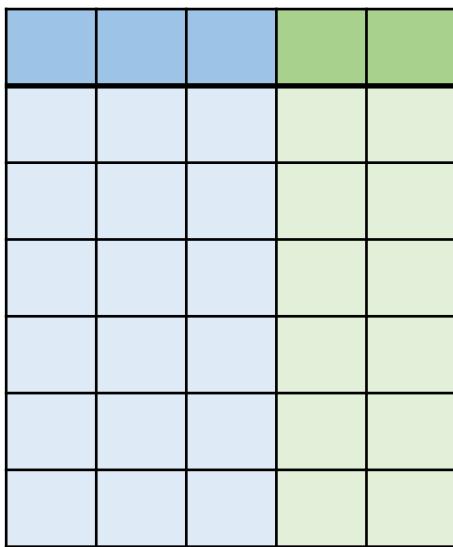


Machine Learning categories

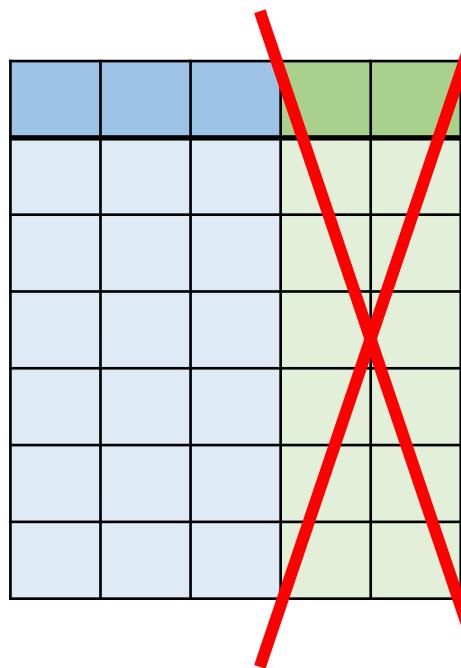


Machine Learning

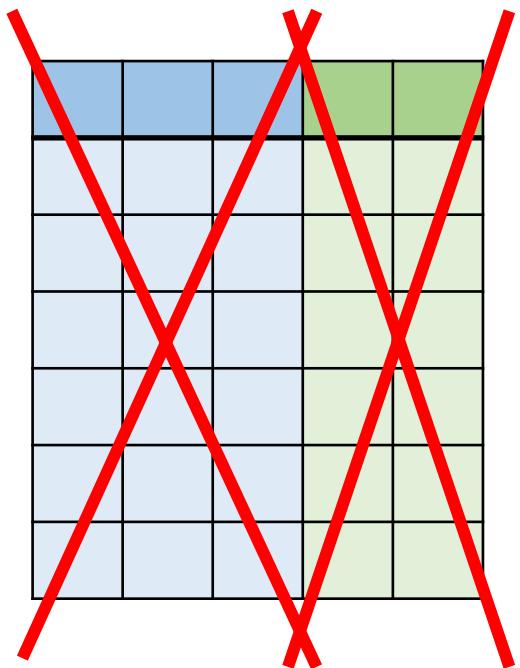
Supervised Learning



Unsupervised Learning



Reinforcement Learning

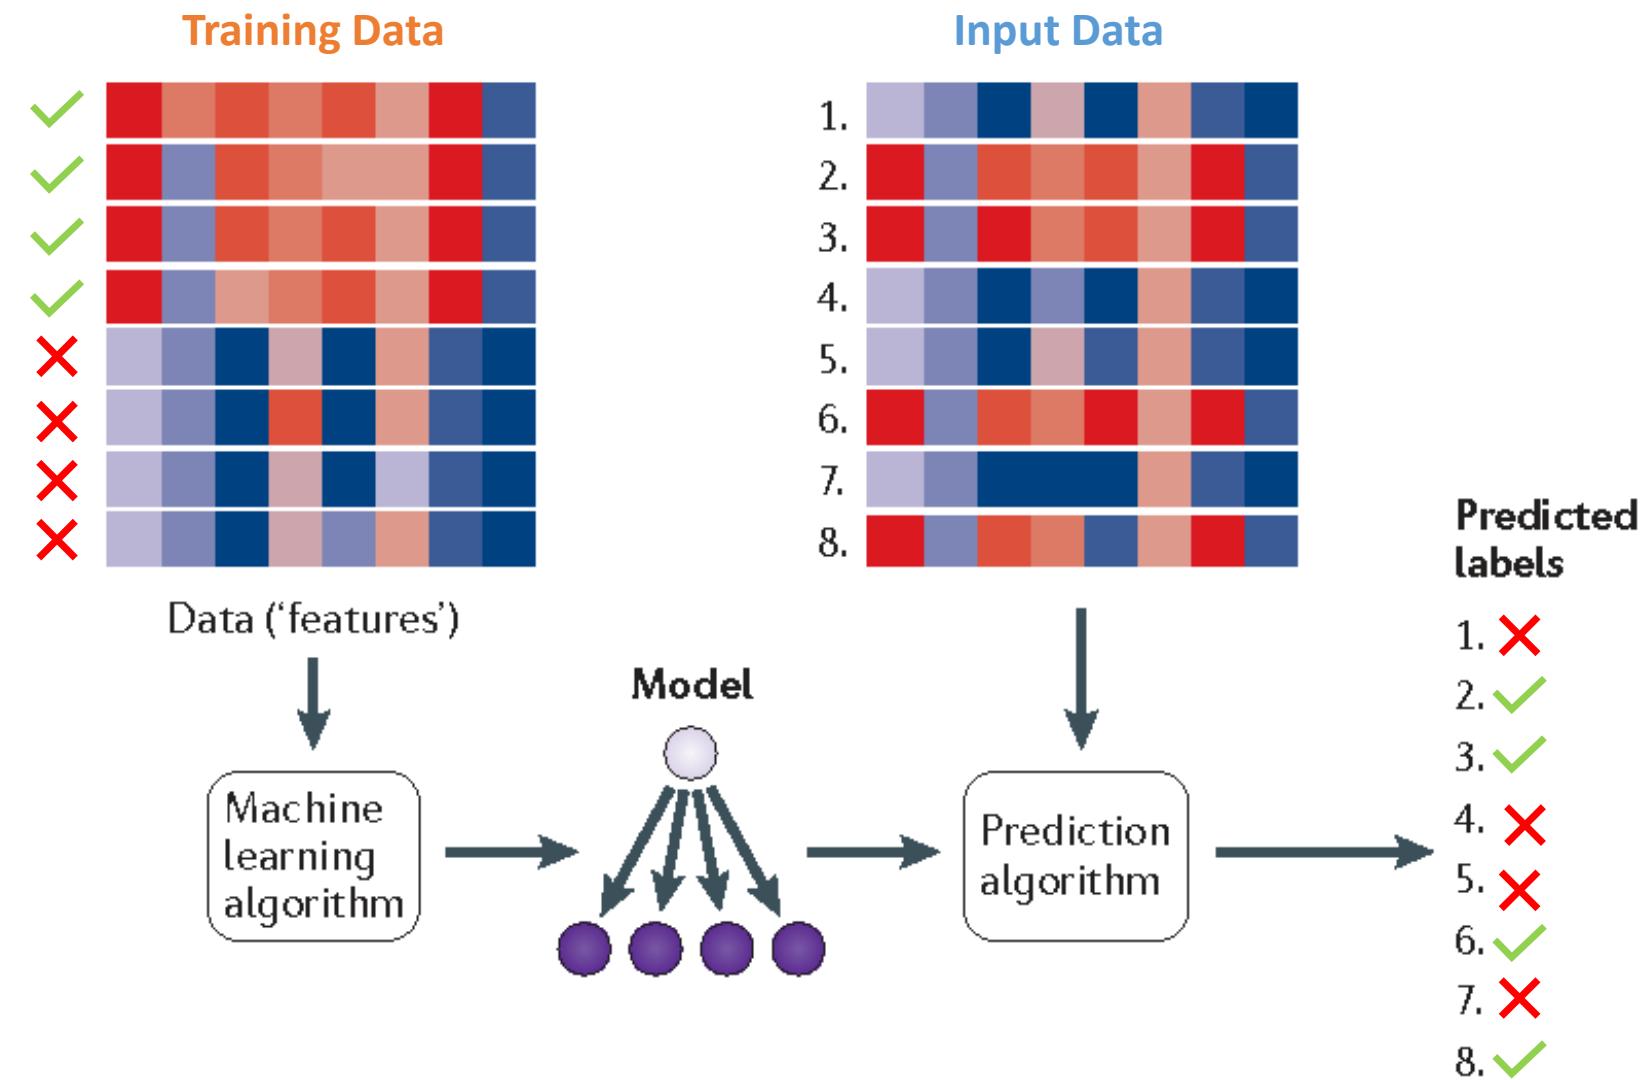


Machine Learning categories

- **Supervised learning:** The computer is presented with **example inputs** and their **desired outputs** (labels, answers, results), given by a "teacher", and the goal is to learn a general rule that maps inputs to outputs.
- **Unsupervised learning:** **No labels** are given to the learning algorithm, leaving it on its own to find structure in its input.
- **Reinforcement learning:** A computer program interacts with a **dynamic environment** in which it must perform a certain goal. The program is provided feedback in terms of **rewards** and punishments as it navigates its problem space.

Today, we focus on
Supervised Learning

Example



Example

airplane														
automobile														
bird														
cat														
deer														
dog														
frog														
horse														
ship														
truck														

What is this ?



Example

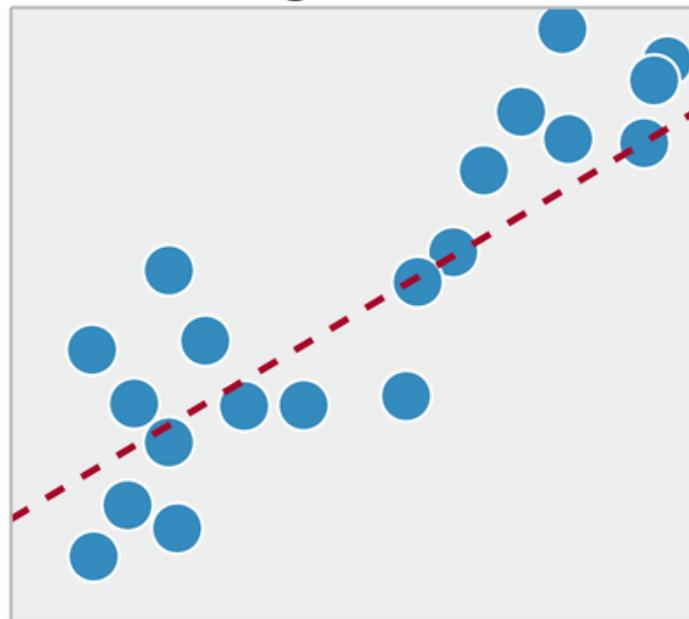
0	0 0 0 0 0 0 0 0 0 0
1	1 1 1 1 1 1 1 1 1 1
2	2 2 2 2 2 2 2 2 2 2
3	3 3 3 3 3 3 3 3 3 3
4	4 4 4 4 4 4 4 4 4 4
5	5 5 5 5 5 5 5 5 5 5
6	6 6 6 6 6 6 6 6 6 6
7	7 7 7 7 7 7 7 7 7 7
8	8 8 8 8 8 8 8 8 8 8
9	9 9 9 9 9 9 9 9 9 9

What is this ?

4

Supervised Learning : What type of problems ?

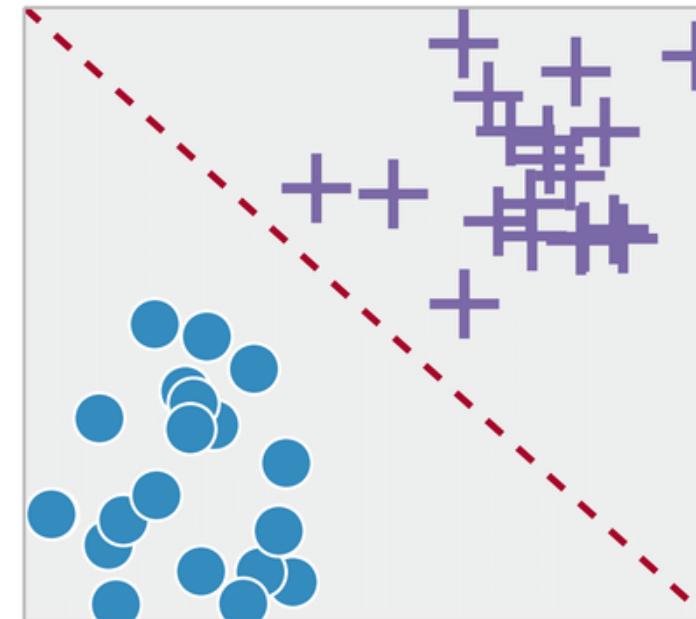
Regression



Predicts continuous values

What will be
the temperature Tomorrow ?

Classification



Predicts categories

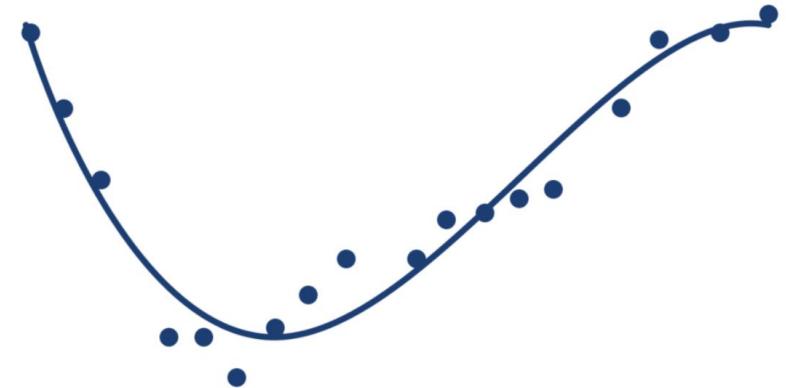
Will it be
Hot or Cold tomorrow ?

Regression problems

- A **regression problem** in ML refers to a type of supervised learning task where the goal is to predict a **continuous output** (numerical value) based on input data (features).
- Examples :
 - What will the temperature be tomorrow?
 - What will bitcoin be worth next week?
 - Estimate house price based on features such as size, number of bedrooms, ..
 - Predict energy consumption of a building based on weather conditions, time of day, and other parameters

Regression algorithms

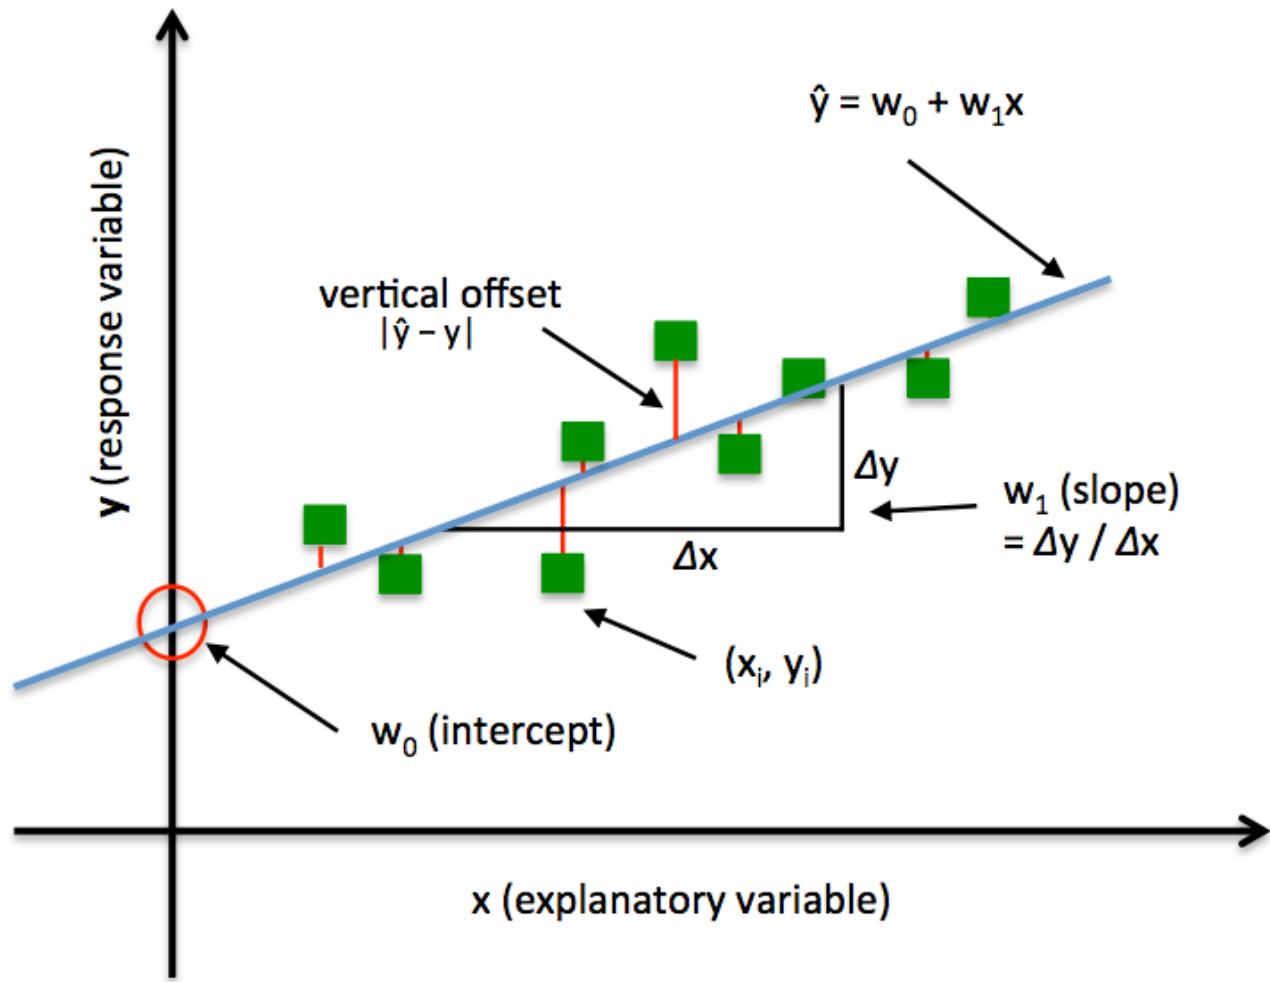
- There are multiple algorithms to model the relation between inputs and outputs :
 - Linear regression, ridge regression, LASSO regression
 - Polynomial regression
 - SVR, KNNR, ...
 - Neural networks !



Linear regression

- Simplest Supervised Learning algorithm, but the most used!
- Predictive analysis
- Estimates the relationship between two independent variables
- The simplest form : $y = c + b*x$
(c: constant, b: regression coefficient)
 - Temperature
 - House price
 - Crime rate
 - Position coordinates
 - Stock price
 - ...

Linear regression : How it works?



$$a = \frac{n \sum_{i=1}^n x_i y_i - \sum_{i=1}^n x_i \sum_{i=1}^n y_i}{n \sum_{i=1}^n x_i^2 - (\sum_{i=1}^n x_i)^2}$$

$$b = \frac{1}{n} \left(\sum_{i=1}^n y_i - a \sum_{i=1}^n x_i \right)$$

Linear Regression in scikit-learn

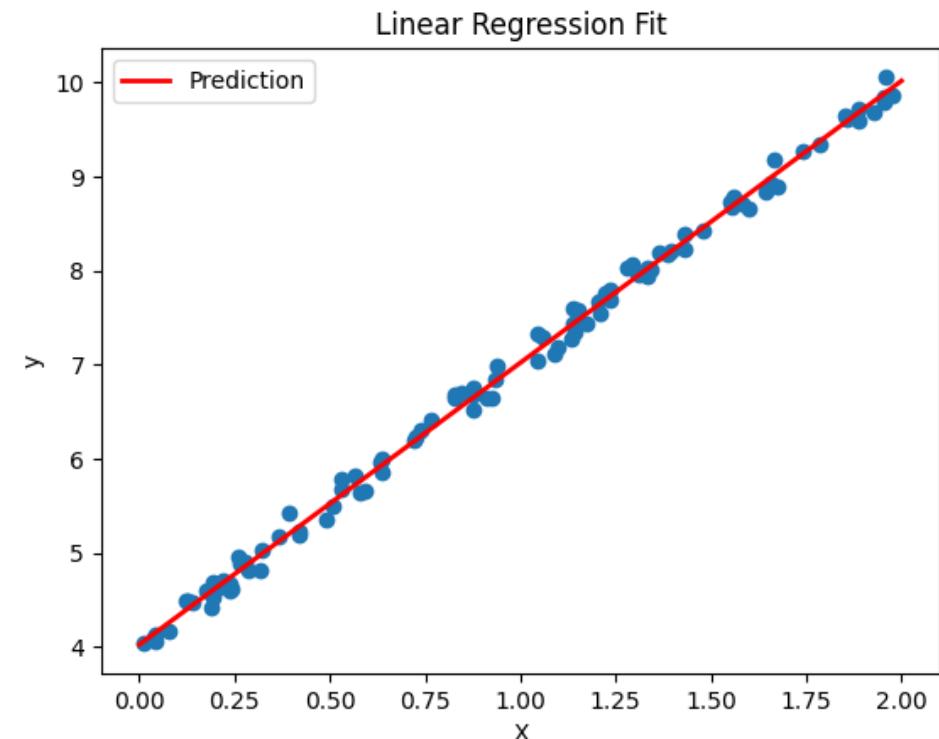
```
import numpy as np
from sklearn.linear_model import LinearRegression

# Generate some data for linear regression
np.random.seed(0)
X = 2 * np.random.rand(100, 1)
y = 4 + 3 * X + 0.1 * np.random.randn(100, 1)

# Perform linear regression using scikit-learn
lr = LinearRegression()
lr.fit(X, y)

# Print the parameters
print("Intercept:", lr.intercept_[0])
print("Slope:", lr.coef_[0][0])

# Use the model to predict the corresponding value for 1.0
print(lr.predict([[1.0]]))
```



Classification problems

- A **classification problem** is a type of supervised learning task where the goal is to predict a **discrete label** (or category) for given input data.
- Examples :
 - Is this mail “spam” or “not spam” ?
 - Is this patient “positive” or “negative” for this disease?
 - What is this character ? (0, 1, ... 9, A, B, C, ...)
 - What object is in this photo ?

Classification Algorithms

- There are multiple algorithms to model the relation between inputs and desired labels :
 - Decision trees
 - Random forests
 - Linear Discriminant Analysis (LDA)
 - K-Nearest Neighbors Regression (KNN)
 - Logistic regression
 - Naive Bayes classification
 - Support Vector Machines
 - Neural networks !

Decision Trees

Definition

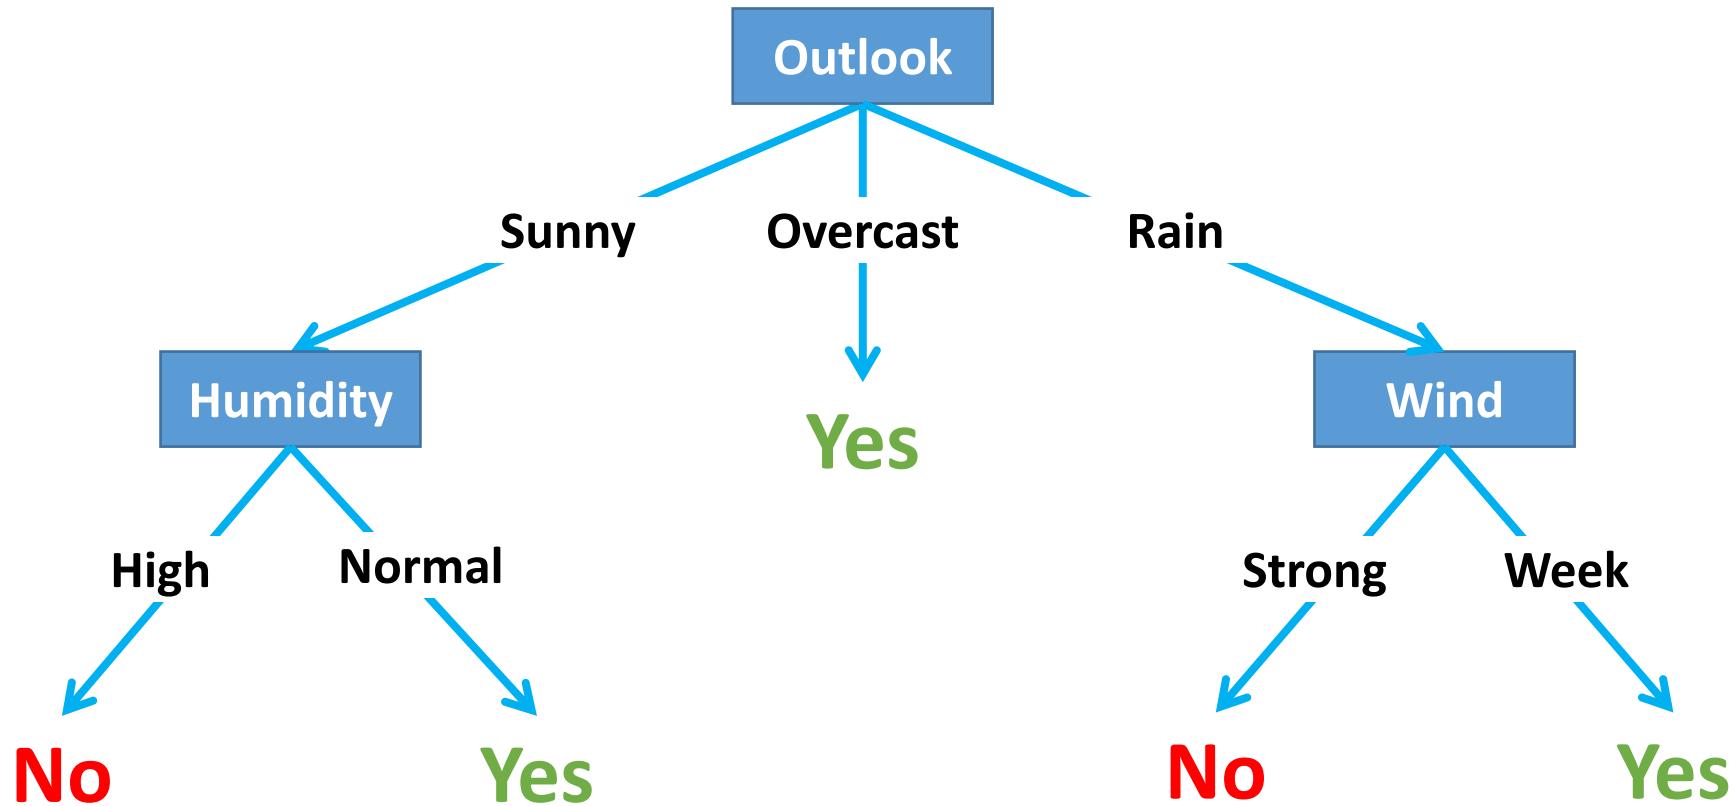
- Decision tree is a classifier in the form of a tree structure
 - **Decision node:** specifies a test on a single attribute
 - **Leaf node:** indicates the value of the target attribute
 - **Arc/edge:** split of one attribute
 - **Path:** a disjunction of test to make the final decision
- Decision trees classify instances or examples by starting at the root of the tree and moving through it until a leaf node.

Training Data Example : The weather problem

Given this data,
is it possible to predict
automatically when
this player will play
tennis or not ??

Day	Outlook	Temperature	Humidity	Wind	PlayTennis
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

Play or not : Decision tree



How to Build a decision tree

- Top-Down Induction of Decision Trees ID3
 1. $A \leftarrow$ the “best” decision attribute for next node
 2. Assign A as decision attribute for node
 3. For each value of A create new descendant
 4. Sort training examples to leaf node according to the attribute value of the branch
 5. **If** all training examples are perfectly classified (same value of target attribute) then stop,
else iterate over new leaf nodes.

Entropy and Information Gain

$$Entropy(S) \equiv -p_{\oplus} \log_2 p_{\oplus} - p_{\ominus} \log_2 p_{\ominus}$$

$$Gain(S, A) \equiv Entropy(S) - \sum_{v \in Values(A)} \frac{|S_v|}{|S|} Entropy(S_v)$$

- Root node: $S = [9+, 5-]$ (all training data: 9 play, 5 no-play)
- Entropy: $H(S) = -(9/14) \log_2(9/14) - (5/14) \log_2(5/14) = 0.94$
- $S_{weak} = [6+, 2-] \Rightarrow H(S_{weak}) = 0.811$
- $S_{strong} = [3+, 3-] \Rightarrow H(S_{strong}) = 1$

$$\begin{aligned} IG(S, \text{wind}) &= H(S) - \frac{|S_{weak}|}{|S|} H(S_{weak}) - \frac{|S_{strong}|}{|S|} H(S_{strong}) \\ &= 0.94 - 8/14 * 0.811 - 6/14 * 1 \\ &= 0.048 \end{aligned}$$

Day	Outlook	Temperature	Humidity	Wind	PlayTennis
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

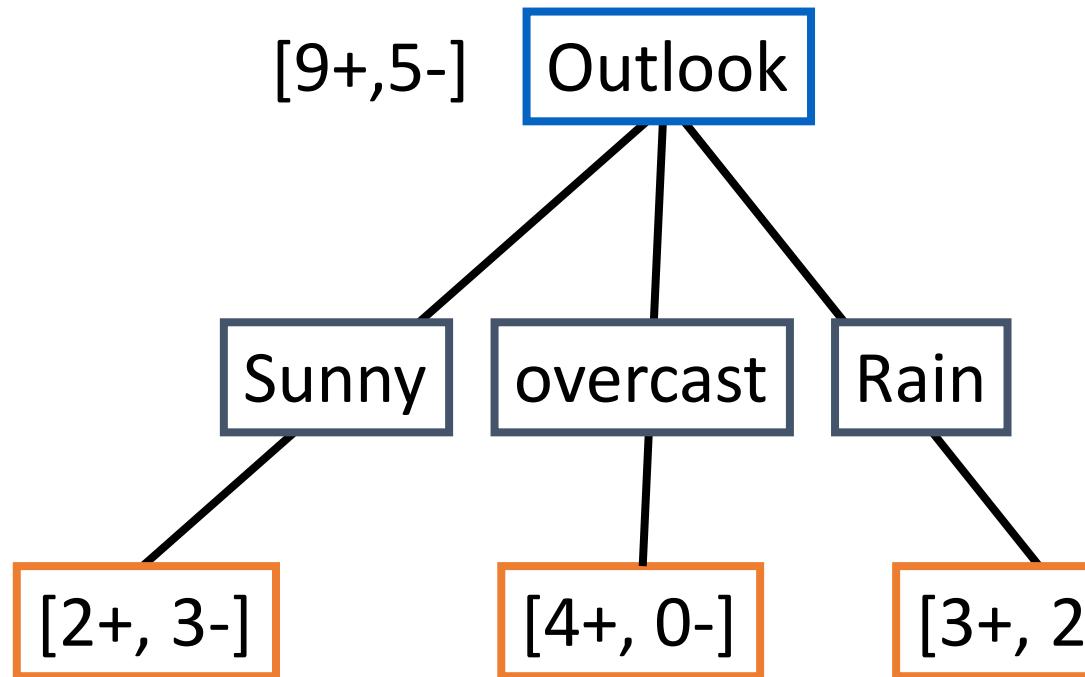
The algorithm computes the IG for all attributes :

- $IG(S, \text{wind}) = 0.048$
- $IG(S, \text{outlook}) = 0.246$
- $IG(S, \text{humidity}) = 0.151$
- $IG(S, \text{temperature}) = 0.029$

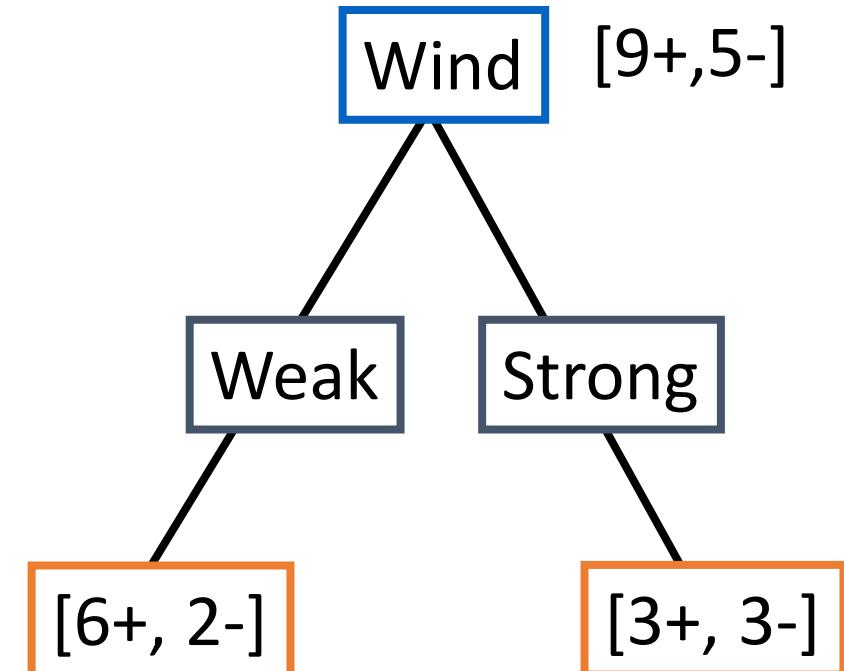
Day	Outlook	Temperature	Humidity	Wind	PlayTennis
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No

Which Attribute is “best”?

It depends on **information gain**

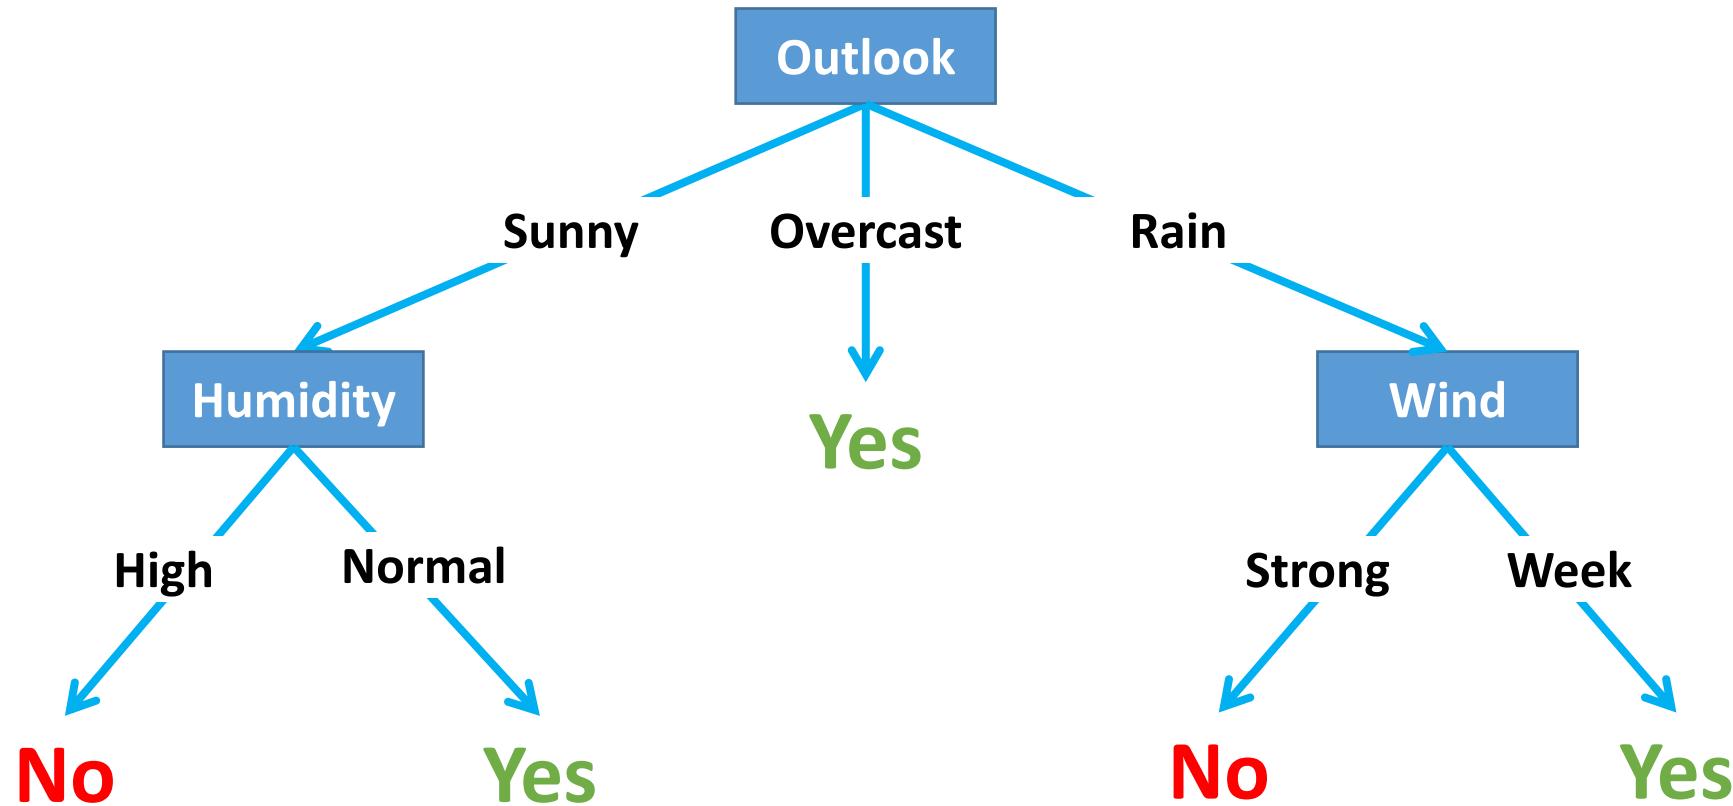


$$IG(S, \text{Outlook}) = 0.246$$



$$IG(S, \text{Wind}) = 0.048$$

Play or not : Decision tree



You are not a data scientist..



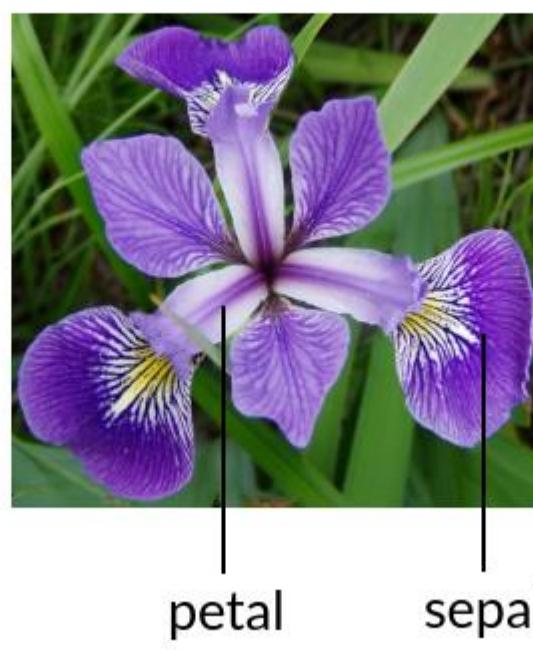
if you don't know this flower

Labwork : iris classification

iris setosa



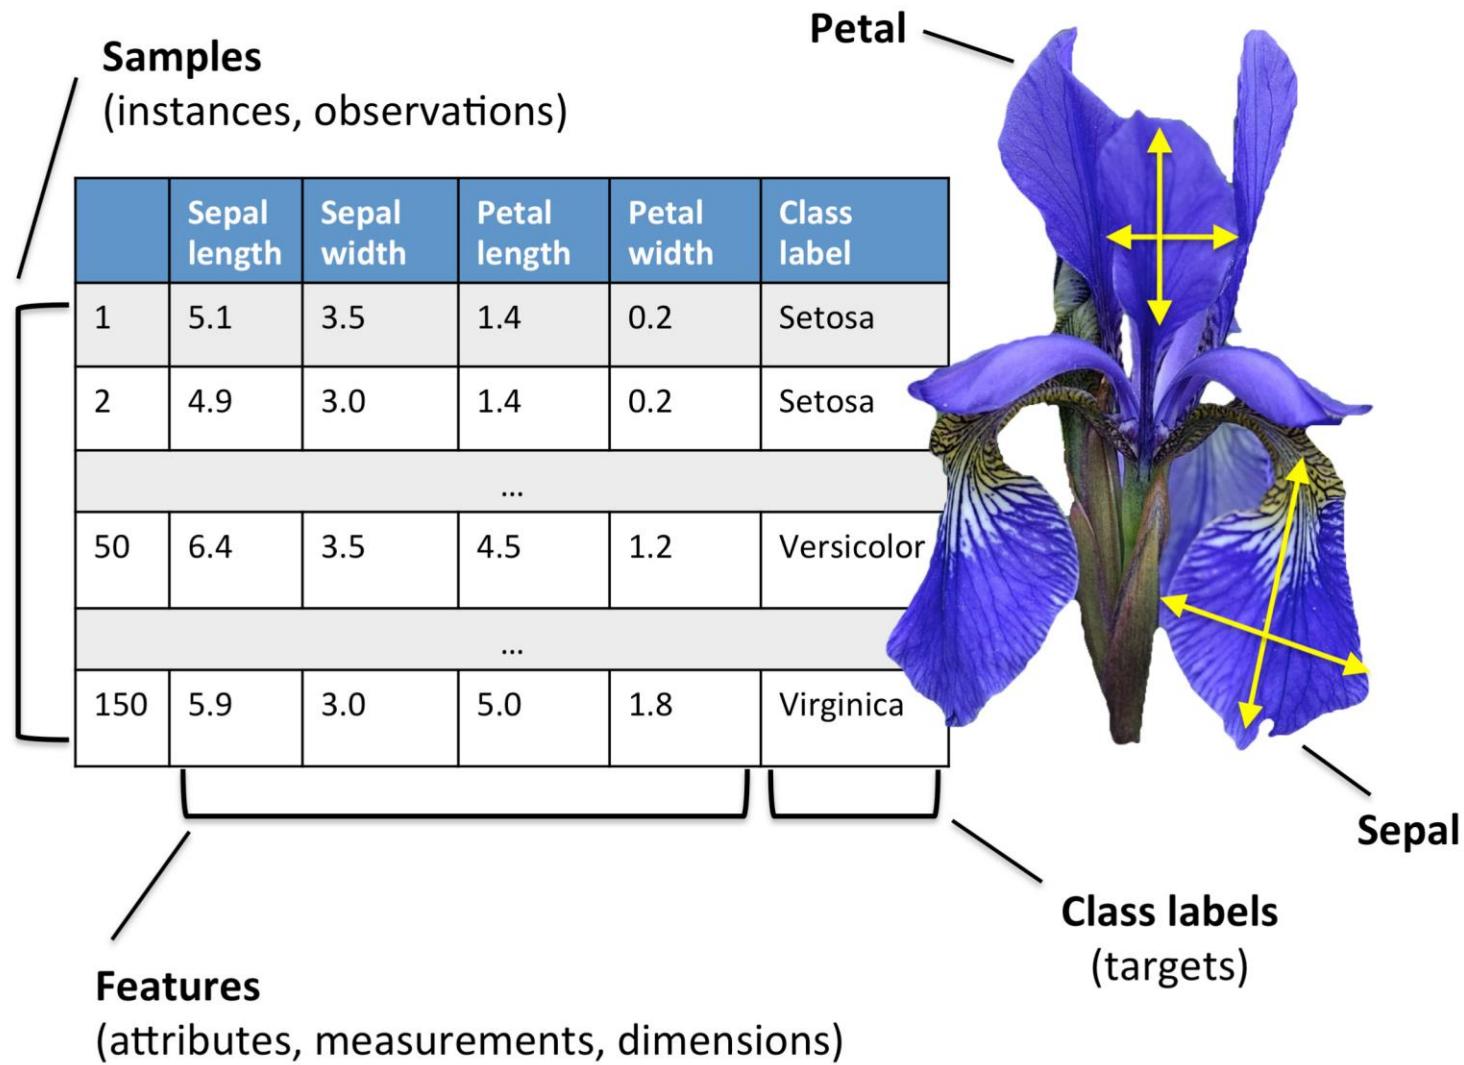
iris versicolor



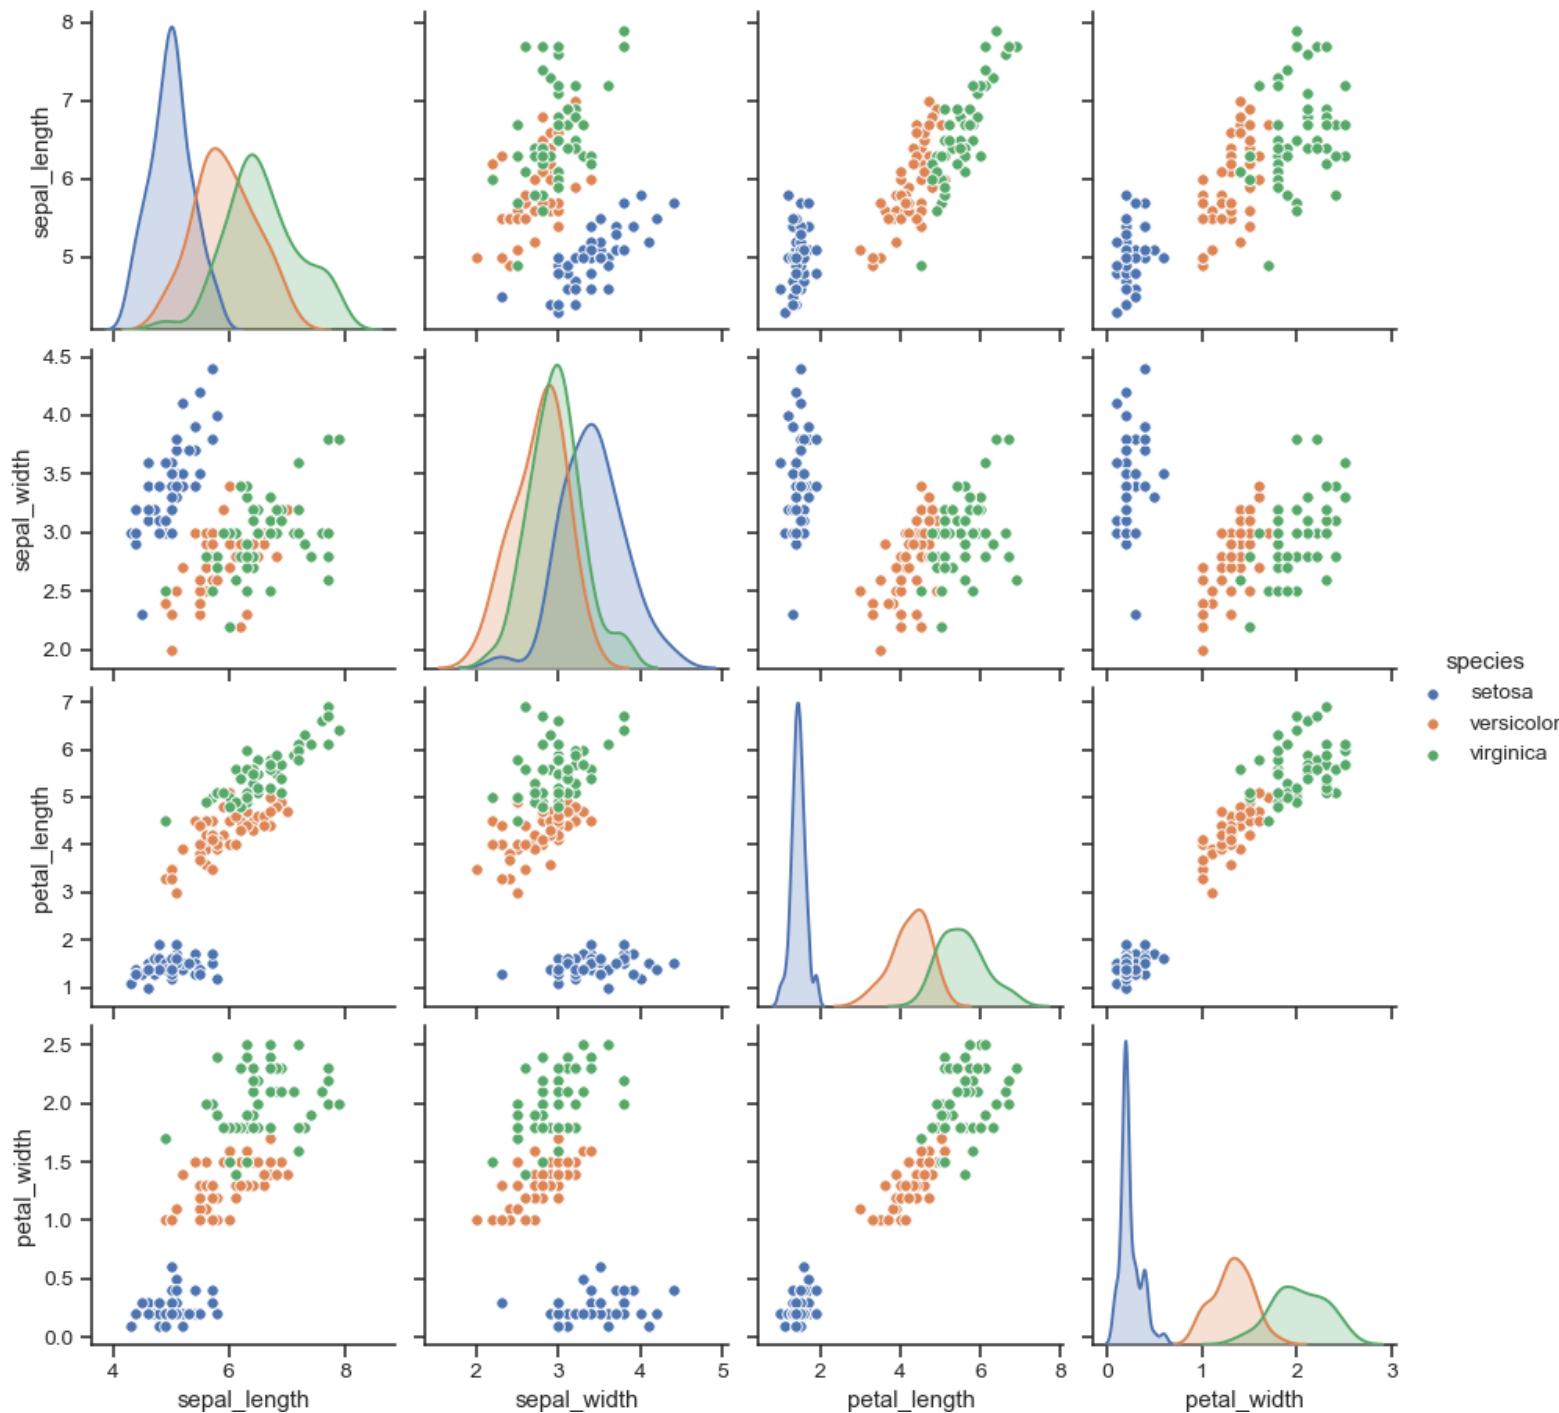
iris virginica



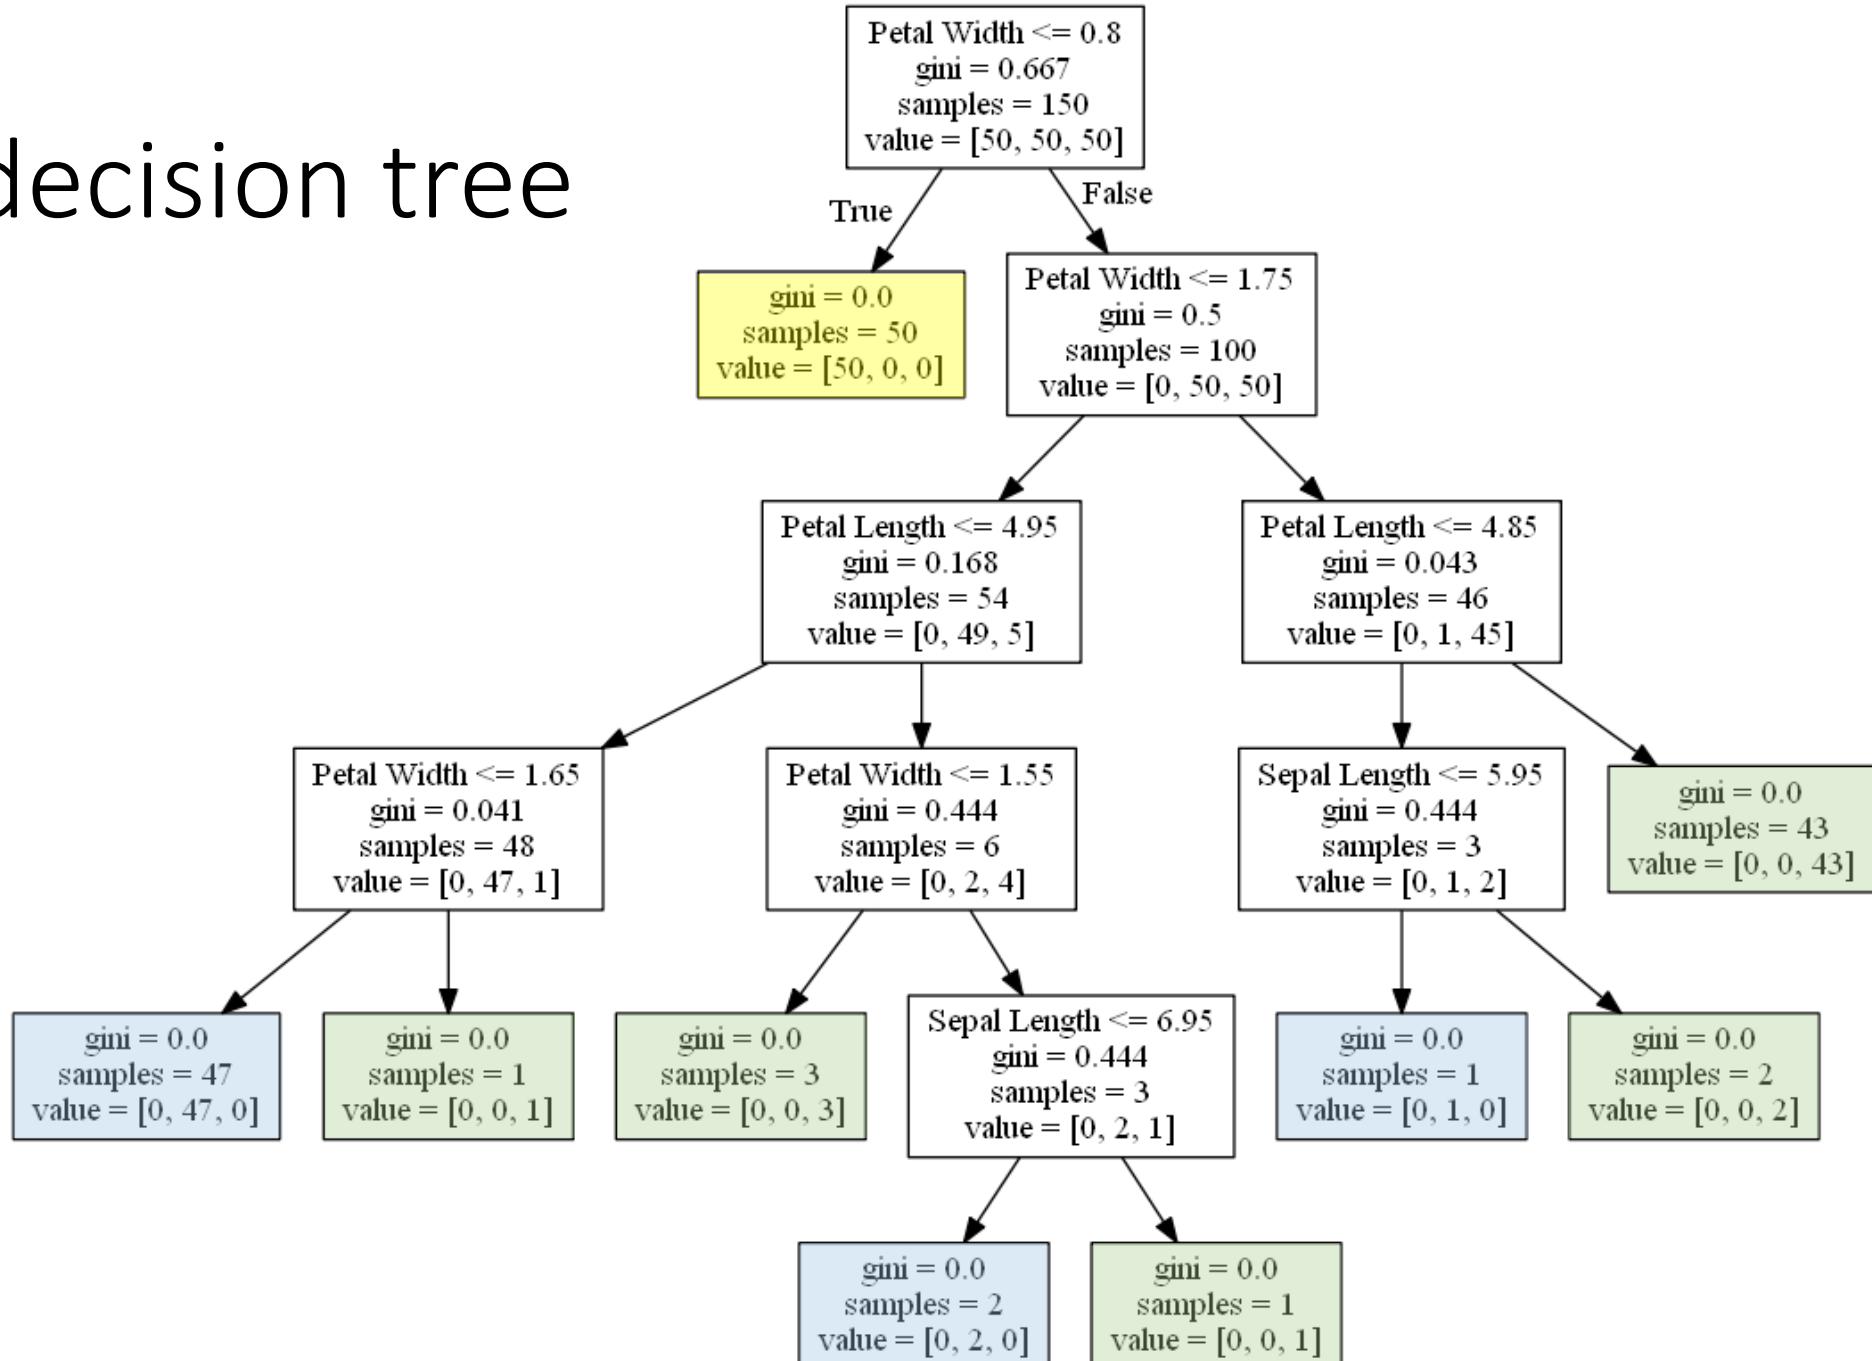
The dataset



Data Visualisation



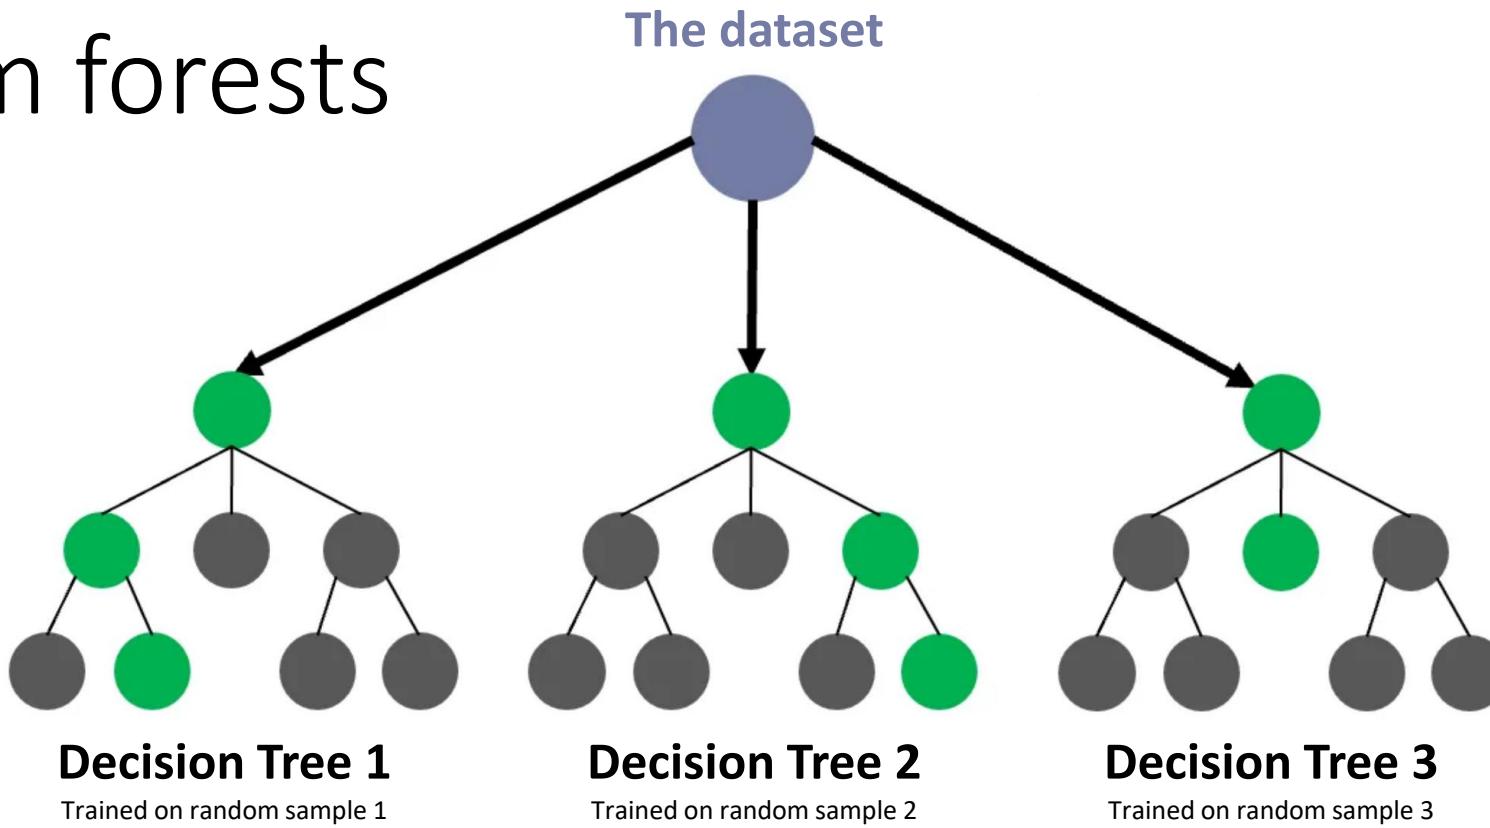
The decision tree



Random forests

Random forests

Training



Predicting

Prediction 1

Prediction 2

Prediction 3

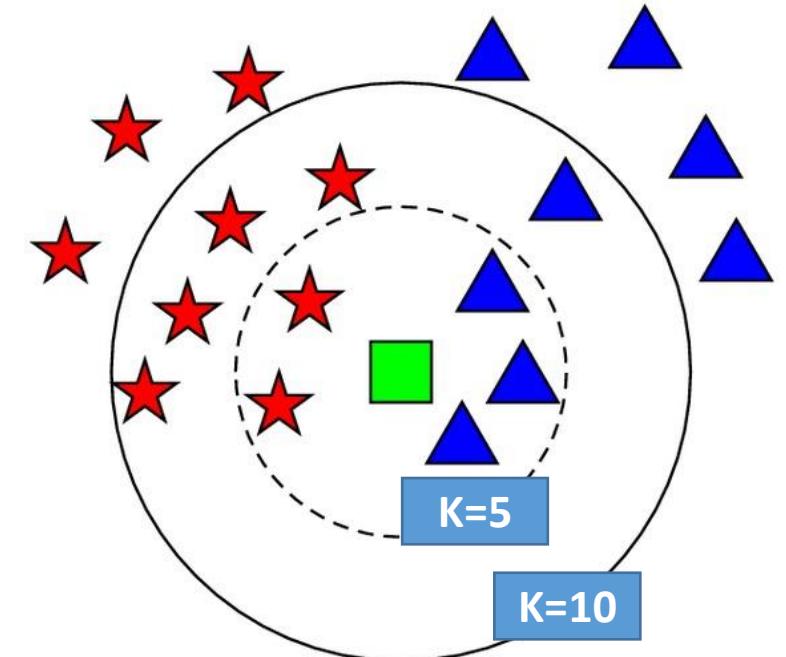
Majority voting / Averaging

Final result

K - Nearest Neighbors

Classification: K-Nearest Neighbors

- Classification of new elements depends on the « k » nearest (closest) elements
- Majority voting / Averaging
- How to choose the value of k ?
- The algorithm creates decision surfaces for faster computation.



Classification Model Performances

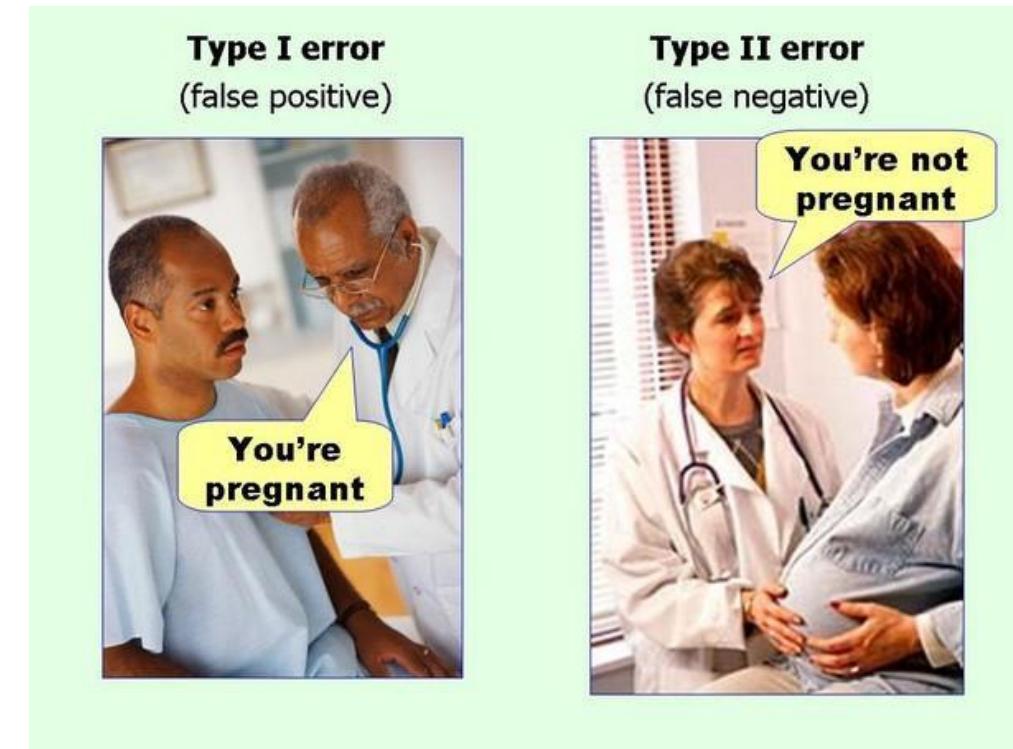
Performances of Classification Models

- The classification model performance is evaluated using **metrics** that measure **how well the model predicts the correct class labels**.
- Metrics include:
 - Accuracy
 - Precision
 - Recall (or sensitivity)
 - F1-score
 - Area Under the Curve (AUC)
 - Receiver Operator Curve (ROC)
 - Likelihood ratios

https://en.wikipedia.org/wiki/F-score#Diagnostic_testing

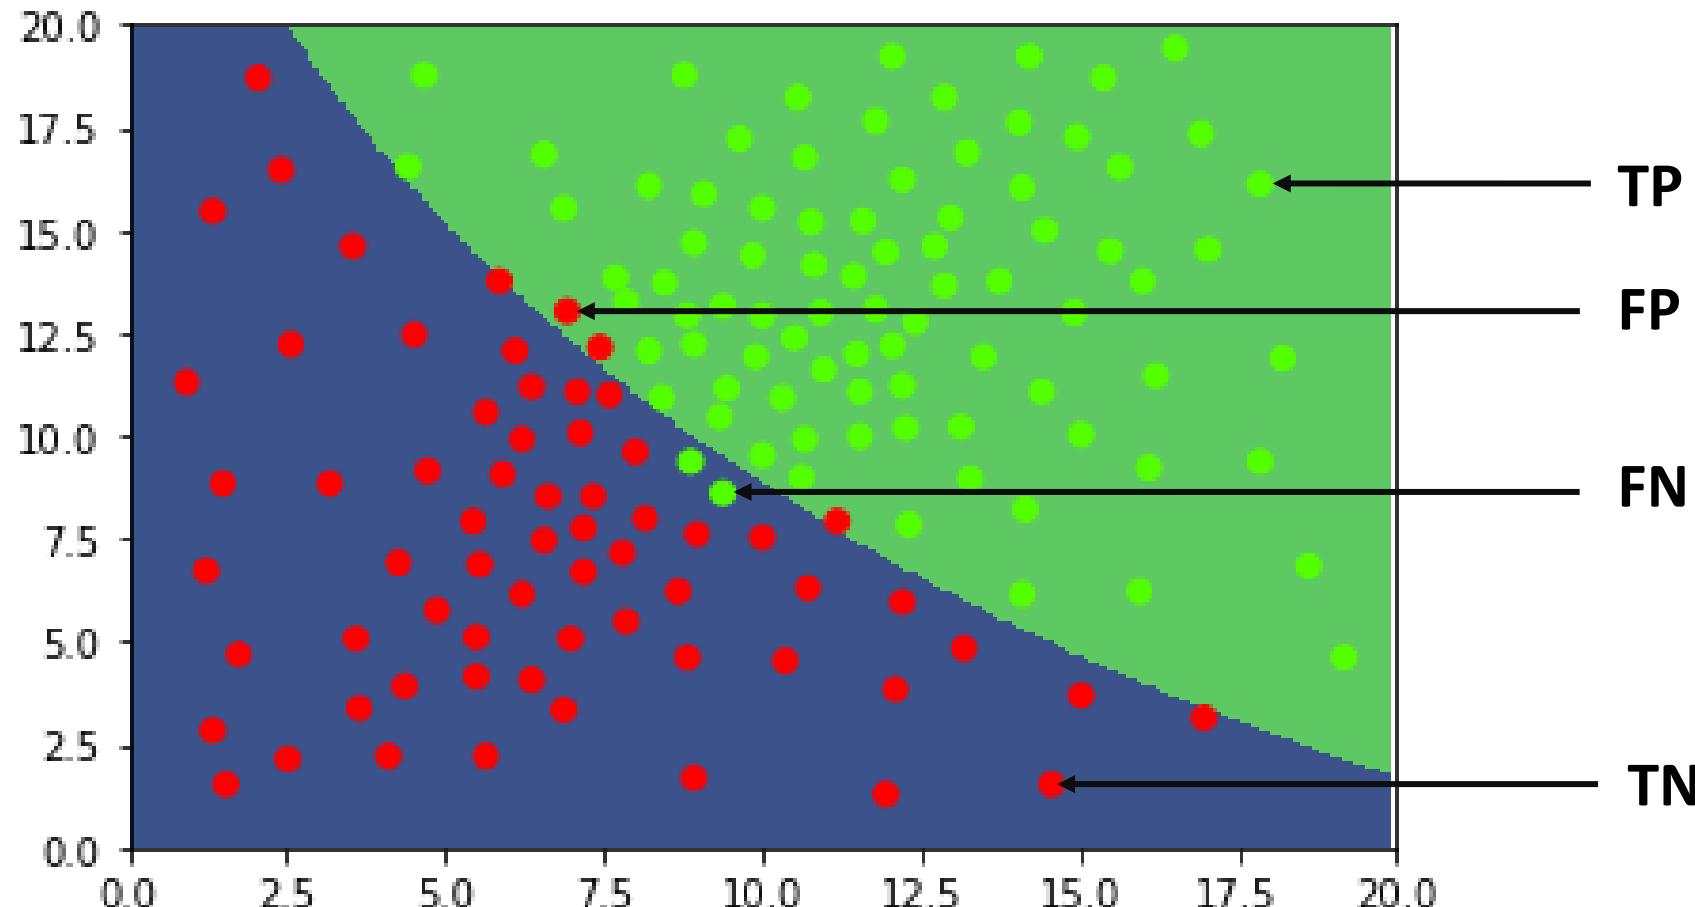
Classification Model Performances

		Predicted class	
		Positive	Negative
Actual class	Positive	TRUE POSITIVE TP	FALSE NEGATIVE FN
	Negative	FALSE POSITIVE FP	TRUE NEGATIVE TN



False predictions

Example: Back to college admission example



Accuracy

- A common metric in classification
- The ratio of good predictions over all predictions

$$\text{Accuracy} = \frac{\text{True Positives} + \text{True Negatives}}{\text{True Positives} + \text{True Negatives} + \text{False Positives} + \text{False Negatives}}$$

```
>>> from sklearn.metrics import accuracy_score
>>> y_pred = [0, 2, 1, 3]
>>> y_true = [0, 1, 2, 3]
>>> accuracy_score(y_true, y_pred)
0.5
```

Precision

- Precision is about the ability of the classifier to don't miss positive cases (not label true negative cases as positives)
- Close to 1 → Good precision

$$\text{Precision} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Positives}}$$

Recall

- The ability of the classifier to find positive cases
- Close to 1 → We found all positive cases

$$\text{Recall} = \frac{\text{True Positives}}{\text{True Positives} + \text{False Negatives}}$$

Example

```
Entrée [42]: from sklearn.metrics import precision_score
              from sklearn.metrics import recall_score
              from sklearn.metrics import accuracy_score
              data_actual = [1, 1, 0, 1, 1, 1, 0, 1, 1, 0]
              data_pred = [1, 0, 1, 1, 1, 0, 0, 1, 0, 0]

              accuracy_score(data_actual, data_pred)
```

Out [42]: 0.6

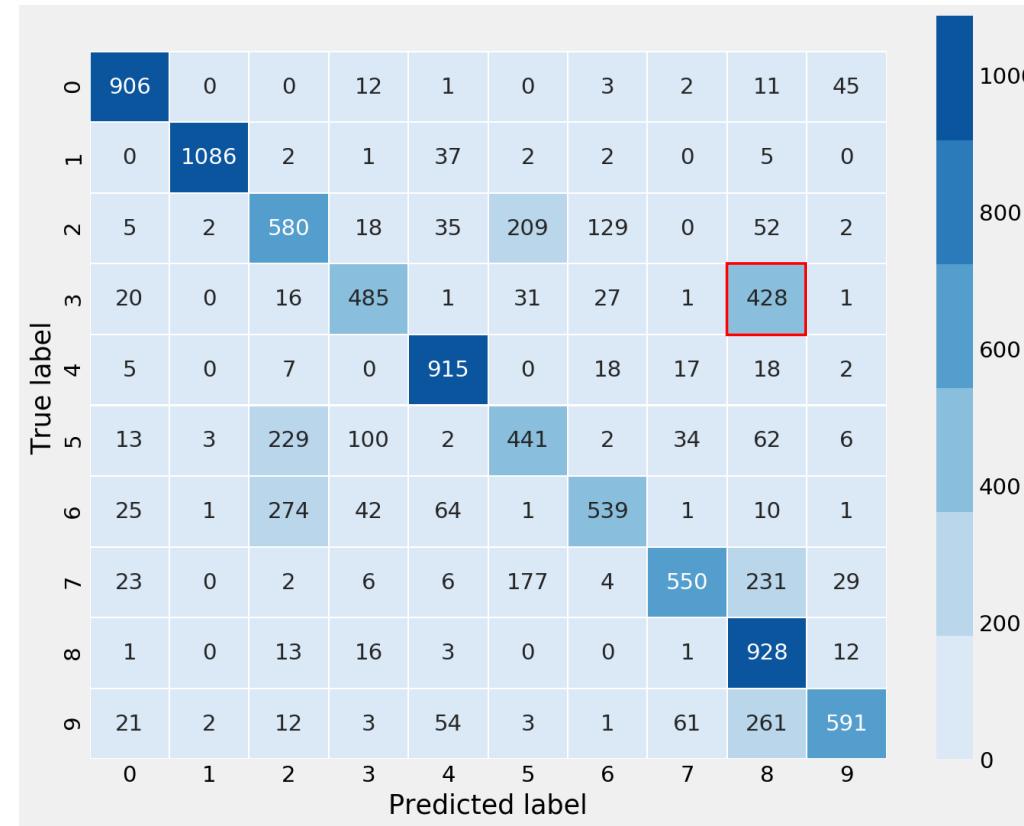
```
Entrée [43]: precision_score(data_actual, data_pred, average='binary')
```

Out [43]: 0.8

```
Entrée [44]: recall_score(data_actual, data_pred, average='binary')
```

Out [44]: 0.5714285714285714

Confusion matrices for multi-label classification

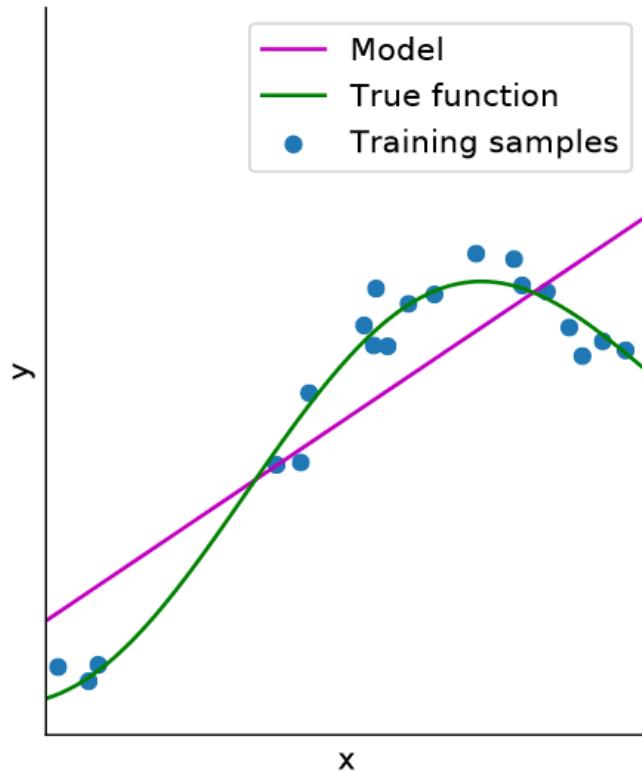


Example of confusion matrix

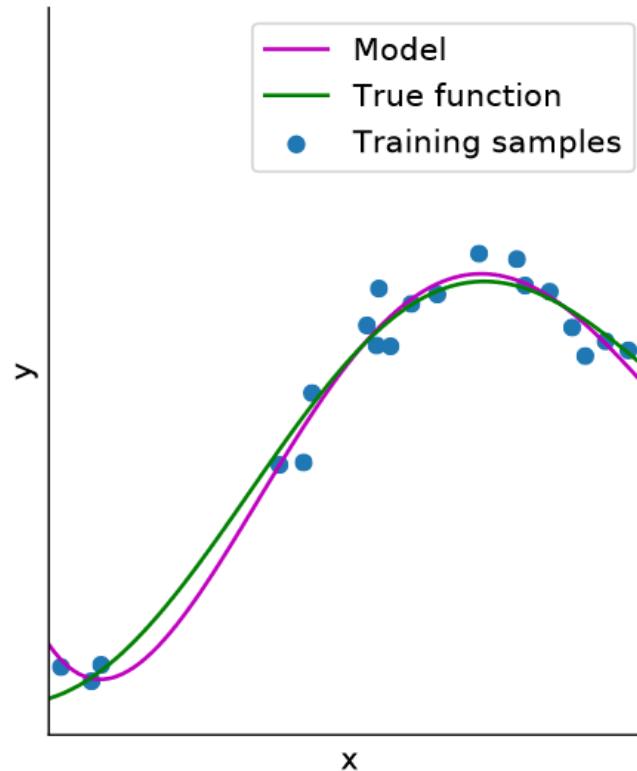
(Observe : A lot of “3”s were classified as “8”s)

Underfitting / Overfitting

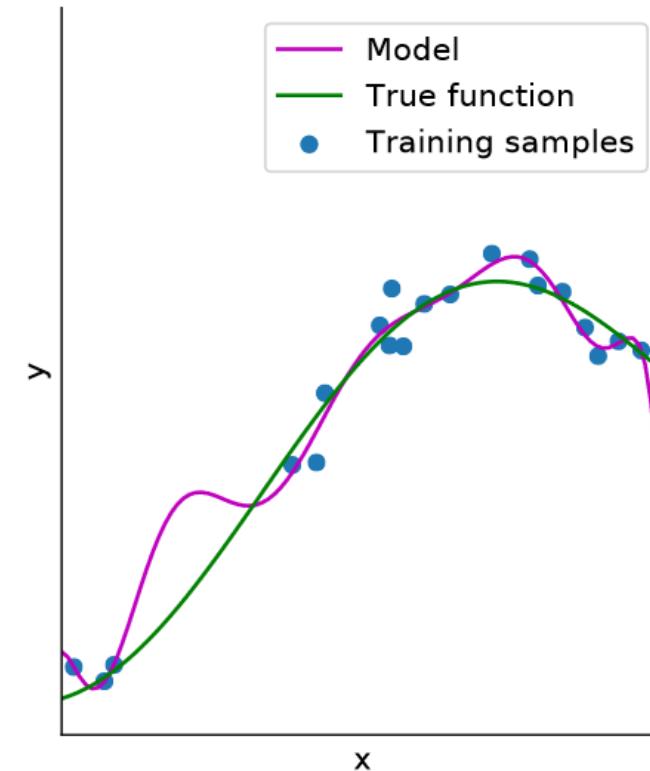
Polynomial of degree 1
for the model



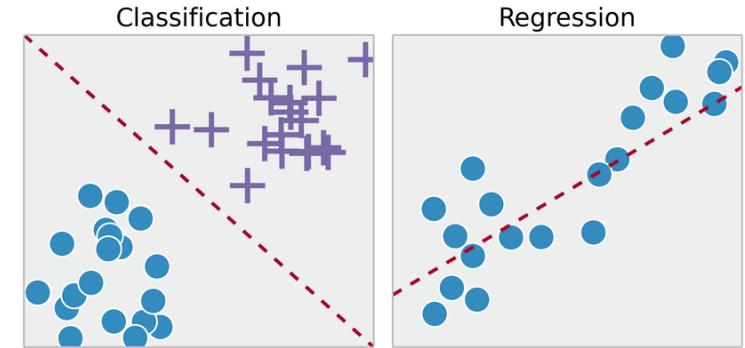
Polynomial of degree 4
for the model



Polynomial of degree 10
for the model



Classification VS Regression



Output type

Classification

Discrete

Objective

Decision boundary

Evaluation

Accuracy, Precision, ...

Regression

Continuous

Best fit line

Sum of squared error